

1      **Supplementary Material: Different antigenic**  
2      **distance metrics generate similar**  
3      **predictions of influenza vaccine response**  
4      **breadth despite low correlation**

5      **Table of contents**

6      1. Reproducibility instructions .....	2
7      2. Extended Methods.....	4
8         2.1 Antigenic distance calculation .....	4
9         2.2 Sequence data sources .....	5
10      2.3 Causal modeling and model formulation .....	6
11      2.4 Model implementation .....	8
12      2.5 Censoring bounds.....	10
13      2.6 Stan implementation.....	11
14      2.7 Posterior marginal effects.....	12
15     3. Supplementary results.....	13
16        3.1 Annual FluZone vaccine formulation.....	13
17        3.2 Annual heterologous strain panel.....	14
18        3.3 Strain names and abbreviations .....	16
19        3.4 Demographic information .....	17
20        3.5 Metric agreement analysis .....	20
21        3.6 Correlation coefficients and CIs .....	22
22        3.7 Antigenic distance evenness and dispersion analysis .....	24
23        3.8 Model diagnostics.....	26
24        3.9 Prior sampling diagnostics.....	27
25        3.10 ELPD Diagnostics.....	28
26        3.11 Pointwise prediction comparisons .....	29
27        3.12 Vaccine-specific Predictions.....	33
28        3.13 Alternative distance metrics .....	34
29     4. References .....	36

30

## 31 1. Reproducibility instructions

32 In order to reproduce our results you should first download the archived repository from  
33 either Zenodo ([DOI: 10.5281/zenodo.15522148](https://doi.org/10.5281/zenodo.15522148)) or clone/download the Git repository  
34 (hosted on GitHub here: <https://github.com/ahgroup/billings-comp-agdist-public>). Note  
35 that if you use different software, software versions, or run the results in a way that differs  
36 from these instructions you may have expected errors or differences between your results  
37 and ours.

38 We ran our analysis pipeline on the University of Georgia’s sapelo2 computing cluster,  
39 which is a distributed computing cluster running CentOS Linux release 7.5 which uses  
40 Slurm to schedule jobs. Our code is written as a targets pipeline (1) and can detect  
41 whether you are running an HPC job in a Slurm environment or not. Notably if you use a  
42 Slurm cluster computing environment which is configured differently from UGA’s sapelo2  
43 environment, you may need to make changes to the Slurm submission script (job.sh) or to  
44 the portion of the \_targets.R script that defines the Slurm jobs. Our code will run on a  
45 local interactive R session as well, and will automatically detect the number of cores  
46 available to use. **Each Bayesian model currently requests 32 cores and we therefore  
47 highly suggest running the main analysis on a cluster computing setup.**

48 Once you’ve downloaded the code, you should open the .Rproj file in Rstudio. Using the R  
49 project file and RStudio is not mandatory, but if you don’t, we assume you know what you  
50 are doing. You can then run the entire analysis pipeline by running targets::tar\_make() in  
51 the console. If you are in an extremely limited computing environment, you can add the  
52 option use\_crew = FALSE to force all targets to execute sequentially (although in this case  
53 you probably do not want to run the bayesian models anyways). You can run the command  
54 targets::tar\_visnetwork(TRUE) to see an interactive graph of our analysis pipeline, and  
55 you can pass a vector of target names to tar\_make() (potentially using tidyselect to  
56 define the vector) to only run those targets. Note that due to differences in OS and file  
57 systems, targets will likely appear outdated for you even though they are up-to-date. We  
58 also do not provide all of our model result files in the GitHub because they are extremely  
59 large (over 100GB) and infeasible to distribute, so if you want to edit or examine the  
60 Bayesian models you will need to rerun the code.

61 You will need the following software requirements to run our code.

- 62     • R version 4.4.1, available from <https://cran.r-project.org>.
- 63     • A working C++ compiler – on a linux cluster this is probably set up for you already,  
64       but on Windows you will need RTools 4.4, which is also available from CRAN. On  
65       MacOS you will need the latest version of the XCode command line tools.
- 66     • The RStudio IDE, available from <https://posit.co/download/rstudio-desktop>.
- 67     • Quarto version 16.40, available from <https://quarto.org>.

- 68     • The `renv` R package, version 1.1.4, available from <https://cran.r-project.org/web/packages/renv/index.html>. It will also attempt to install itself the  
69       first time you open our R project.  
70  
71     • Multiple system dependencies, including CMake. On Windows/MacOS these are  
72       provided by RTools or XCode respectively. On any type of linux there may be  
73       additional system requirements you will need to download. Your system should  
74       prompt you about this.  
75     • The packages specified in the file `renv.lock`, which can be installed as explained in  
76       the next section.

77 With the software installed, follow these instructions to reproduce our results.

- 78     1. Open the `billings-comp-agdist-public.Rproj` file in Rstudio.  
79     2. Once `renv` initializes, run the command `renv::restore()` in the console to begin  
80       installing the required packages. If you have issues at this stage you can also install  
81       the dependencies manually, but if you do not use `renv` or you use different package  
82       versions than we did, our code might not work for you.  
83     3. If you want to run any steps that involve Bayesian models, you need to install  
84       `cmdstan` following the `cmdstanr` quick start guide at this location: <https://mc-stan.org/cmdstanr/articles/cmdstanr.html>. We used `cmdstan` version 2.36.0 for  
85       this project.  
86     4. Now you can run our pipeline by running `targets::tar_make()` in the console. If  
87       you are new to `targets` and want to learn more about how the pipeline works, we  
88       recommend reading the `targets` manual which can be found here:  
89       <https://books.ropensci.org/targets/>.

91 **Again we note that our code is computationally intensive and we ran it on a distributed  
92 computing cluster. It still took multiple days to run, even running many operations in  
93 parallel with many cores each.**

94

95 **2. Extended Methods**

96 **2.1 Antigenic distance calculation**

97 We calculated four different antigenic distance metrics for our study. In this section, we  
98 walk through how each method is calculated. Note that we only considered pairwise  
99 distances between strains of the same subtype. So we only computed distances between  
100 two A(H1N1) strains, between two A(H3N2) strains, or between two influenza B strains, we  
101 did not compute distances between A(H1N1) and A(H3N2) strains or between any A and B  
102 strains. However, since the two B lineages are quite similar and our panel included pre-  
103 divergence influenza B strains, we performed pairwise comparisons of all influenza B  
104 strains.

105 **Temporal distance** is the absolute value of the difference in the years of isolation between  
106 the two strains. For example, the difference between A/H1N1/California/09 and  
107 A/H1N1/Michigan/15 would be  $|2015 - 2009| = 6$ . Notably, in our study, we did not have  
108 any examples where the assay strain was isolated later than the vaccine strain, so taking  
109 the absolute value is not necessary, but we wanted to avoid confusion about our  
110 definitions. Future studies that collect such data might prefer either “backwards” or  
111 “forwards” temporal distances, and we can’t comment on that here.

112 **Dominant *p*-Epitope distance** is the maximum length-normalized Hamming distance  
113 across the five major epitope sites on the HA head. After aligning the HA amino acid  
114 sequences for all of the strains, we removed the signal peptides from the sequences and  
115 used the previously identified epitope site locations for influenza A (2) and influenza B (3).  
116 Working pairwise with the sequences, we concatenated the residues for each epitope and  
117 calculated the Hamming distance between each epitope, and we divided the Hamming  
118 distance for a given epitope by the number of residues in that epitope. Then the *p*-Epitope  
119 distance for that pair of strains was the maximum of those epitope-wise distances.

120 **Grantham’s distance** is a weighted distance based on biochemical properties that  
121 considers how different two differing residues at the same position are. We used  
122 Grantham’s substitution matrix (4) to assign a value to each residue site between two  
123 sequences, based on the transition between amino acids. More different transitions are  
124 given higher weights. Then, for each pair of sequences, we sum the weights for that pair  
125 and divide by the length of the sequence.

126 Finally, **cartographic distance** is the Euclidean distance between strains on antigenic  
127 cartography map. We built our cartographic maps from the combined table of post-  
128 vaccination titer data in our study, treating all person-years as independent occurrences  
129 (there is no clear meaning for repeat measurements in a dimension reduction analysis).  
130 We used Racmacs, which implements metric multidimensional scaling, to create and  
131 optimize the cartographic map (5). All of our maps were two dimensional, and we selected  
132 the best fitting map from 25 distinct Racmacs runs with random initializations, where each  
133 initialization was allowed to perform up to 100 L-BFGS optimization runs to relax the initial

134 MDS cartography. Multiple optimization runs are necessary because different initial  
135 conditions can lead to different maps (6). Combining multiple runs by applying a method  
136 like generalized Procrustes analysis is theoretically possible (simple averaging won't work  
137 because rotation and scaling need to be taken into account) but has not yet been studied  
138 or published so we instead chose the one overall best run. We did not perform dimensional  
139 analysis to choose 2D maps, we chose them for ease of interpretation and based on  
140 previous convention.

141 For our models, we only considered the antigenic distance between the assay strain and  
142 the vaccine strain of the same subtype for a given HAI assay. Some of the assay strains  
143 used were influenza B strains isolated before the Victoria/Yamagata lineage divergence.  
144 Because our main question was about the antigenic distance, we compared pre-  
145 divergence B strains to both the Yamagata and Victoria vaccine strains in our analyses. To  
146 facilitate fair comparisons across subtypes and antigenic distance metrics, we min-max  
147 normalized the antigenic distance measurements within each combination of influenza  
148 season, subtype, and metric. After normalization, the antigenic distance for homologous  
149 measurements was set to 0, and the antigenic distance for the most different assay strain  
150 used in a given season was set to 1, with all other antigenic distance values falling in this  
151 interval.

## 152 2.2 Sequence data sources

153 We retrieved HA sequences for each strain from either the U.S. National Center for  
154 Biotechnology Information (NCBI)'s GenBank database (7,8), the UniProt dataset (9), or  
155 GISAID's EpiFlu database (10,11). The attribution and accession numbers for each strain  
156 are listed in [Table 1](#).

157 Most of the sequences we used from GenBank and UniProt are not associated with  
158 particular publications and are only able to be referenced via their accession numbers. The  
159 following sequences have formal references: AAD17229.1 (12); AAA67338.1 (13);  
160 AAP34324.1 (14); ADE28750.1 (15); ACP41953.1 (16); ABQ97200.1 (17); AAA62338.1 (18);  
161 AIW60702.1 (19); P03460 and P03461 (20); and P12443 (21).

162 The sequences we used from GISAID are accessible via GISAID Identifier  
163 EPI\_SET\_250609vz and DOI <https://doi.org/10.55876/gis8.250609vz>. To view the  
164 contributors of each individual sequence with details such as accession number, Virus  
165 name, Collection date, Originating Lab and Submitting Lab and the list of Authors, visit  
166 [10.55876/gis8.250609vz](https://doi.org/10.55876/gis8.250609vz).

*Table 1: Accession number and source for each HA sequence used in our analysis.*

Strain Name	Source	Accession #
A/H1N1/South Carolina/1/1918	GenBank	AAD17229.1
A/H1N1/Puerto Rico/8/1934	GenBank	AGU93019.1
A/H1N1/Weiss/1943	GenBank	ABD79101.1
A/H1N1/Fort Monmouth/1/1947	GenBank	AAA67338.1
A/H1N1/Denver/1957	GenBank	ABD15258.1
A/H1N1/New Jersey/8/1976	GenBank	AGB51356.1
A/H1N1/Ussr/90/1977	GenBank	ABD95350.1

A/H1N1/Brazil/11/1978	GenBank	ABO38065.1
A/H1N1/California/10/1978	GenBank	ABP49338.1
A/H1N1/Chile/1/1983	GenBank	ABO38340.1
A/H1N1/Singapore/6/1986	GenBank	ABO38395.1
A/H1N1/Texas/36/1991	GenBank	ACF41933.1
A/H1N1/Beijing/262/1995	GenBank	ACF41867.1
A/H1N1/New Caledonia/20/1999	GenBank	AAP34324.1
A/H1N1/Solomon Islands/3/2006	GenBank	ABU99109.1
A/H1N1/Brisbane/59/2007	GenBank	ADE28750.1
A/H1N1/California/07/2009	GenBank	ACP41953.1
A/H1N1/Michigan 45/2015	GenBank	AMV49034.1
A/H1N1/Brisbane/02/2018	GISAID	EPI1415369
A/H1N1/Guangdong-Maonan/SWL1536/2019	GISAID	EPI3133357
A/H1N1/Victoria/2570/2019	GenBank	WEY08940.1
A/H3N2/Hong Kong/8/1968	GenBank	ABQ97200.1
A/H3N2/Port Chalmers/1/1973	GenBank	ABE12532.1
A/H3N2/Texas/1/1977	GenBank	AFM68965.1
A/H3N2/Mississippi/1/1985	GenBank	AAA62338.1
A/H3N2/Sichuan/2/1987	GenBank	AFG72085.1
A/H3N2/Shandong/9/1993	GenBank	AFH00285.1
A/H3N2/Nanchang/933/1995	GenBank	AFG72625.1
A/H3N2/Sydney/5/1997	GenBank	ACO95259.1
A/H3N2/Panama/2007/1999	GenBank	ABF21273.1
A/H3N2/Fujian/411/2002	GenBank	AFG72823.1
A/H3N2/New York/55/2004	GenBank	ACF41900.1
A/H3N2/Wisconsin/67/2005	GenBank	AHG96791.1
A/H3N2/Brisbane/10/2007	GenBank	AIW60702.1
A/H3N2/Uruguay/716/2007	GenBank	ACD47213.1
A/H3N2/Perth/16/2009	GenBank	ACS71642.1
A/H3N2/Victoria/361/2011	GenBank	AIU46088.1
A/H3N2/Texas/50/2012	GenBank	AGL07159.1
A/H3N2/Switzerland/9715293/2013	GISAID	EPI530687
A/H3N2/Hong Kong/4801/2014	GISAID	EPI1834581
A/H3N2/Singapore/infimh-16-0019/2016	GISAID	EPI780183
A/H3N2/Kansas/14/2017	GenBank	AVG71503.1
A/H3N2/South Australia/34/2019	GISAID	EPI1387331
A/H3N2/Hong Kong/2671/2019	GenBank	WMW30924.1
A/H3N2/Tasmania/503/2020	GenBank	WMW30850.1
A/H3N2/Darwin/9/2021	GenBank	WND60806.1
B/Lee/1940	UniProt	P03460
B/Maryland/1959	UniProt	P03461
B/Singapore/1964	UniProt	P12443
B/Victoria/02/1987	UniProt	A4D5N9
B/Hong Kong/330/2001	GenBank	ABL77178.1
B/Malaysia/27127/2004	GenBank	AFJ80733.1
B/Victoria/326/2006	GenBank	AGX18732.1
B/Brisbane/60/2008	GenBank	AFH57909.1
B/Colorado/06/2017	GenBank	ASK81305.1
B/Washington/02/2019	GenBank	WIM08940.1
B/Michigan/01/2021	GenBank	WMW30908.1
B/Austria/1359417/2021	GISAID	EPI1868375
B/Yamagata/16/1988	GenBank	ABL77255.1
B/Harbin/7/1994	GenBank	ACR15721.1
B/Sichuan/379/1999	GISAID	EPI2085837
B/Florida/4/2006	GenBank	ACA33493.1
B/Wisconsin/01/2010	GenBank	AET22057.1
B/Texas/06/2011	GenBank	AGI64713.1
B/Massachusetts/02/2012	GenBank	AGL06036.1
B/Phuket/3073/2013	GISAID	EPI3555941

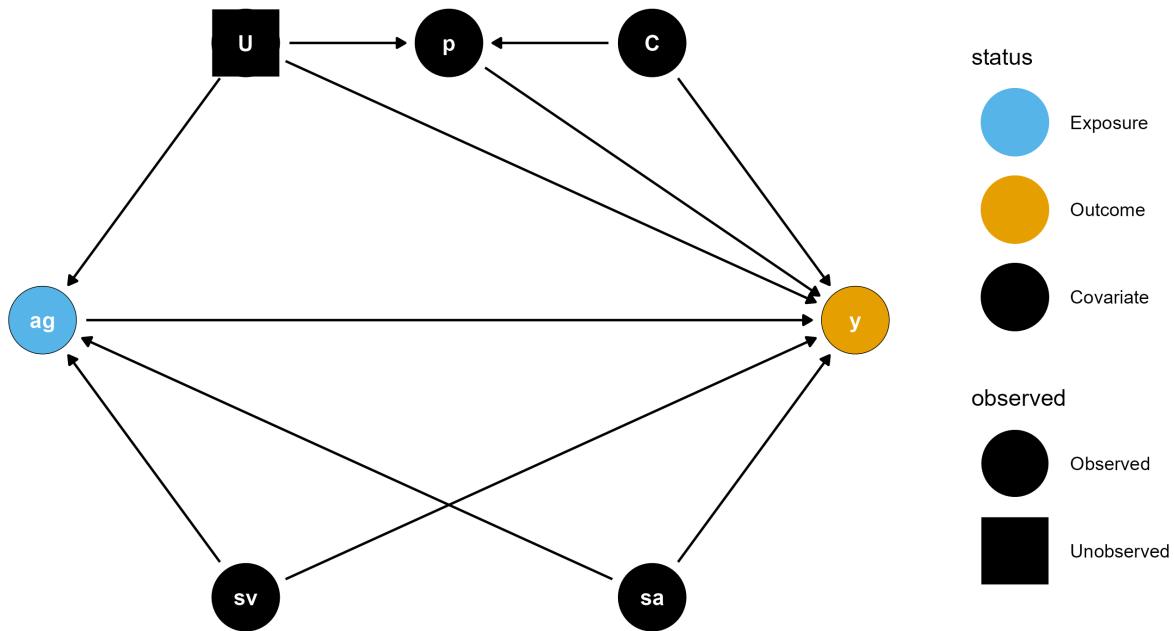
## 167 2.3 Causal modeling and model formulation

168 While we do not claim that our estimates are causal, we employed a graphical causal  
 169 model to formulate our statistical models. While all statistical models are a mix between  
 170 practicality and the best possible model, we hope that by formalizing our thinking, our  
 171 models will be robust and correctly answer our research questions.

172 Our original dataset contained one record per HAI assay, indicating the individual, season,  
 173 study site, time point (pre- or post-vaccination), vaccine dose, and assay strain for each

174 record. The data also included the following demographic variables: age, birth year, sex  
175 assigned at birth, and reported race/ethnicity. The study also provided a list of vaccine  
176 strains for each formulation (see the section on vaccine formulation for a complete list).  
177 Note that we only analyzed standard dose vaccine recipients in our analysis, so we do not  
178 discuss the vaccine dose further.

179 We built a causal model for the effect of antigenic distance as a directed acyclic graph  
180 (DAG). We include the following variables in our causal model:  $U$ , unobserved confounders  
181 that could be partially explained by nuisance variation, but are not directly explained in our  
182 model;  $p$ , the pre-vaccination titer;  $C$ , the set of individual covariates which could  
183 potentially impact both pre-vaccination and post-vaccination titers;  $sv$ , the vaccine strain  
184 (for a given subtype); and  $sa$ , the assay strain for a particular HAI assay. We considered  $C$   
185 to represent all observed demographic features reported in our study, including age, birth  
186 year, sex assigned at birth, and race/ethnicity. Other sources of nuisance variation,  
187 including the study site an individual reported to and other sources of individual variation  
188 which we have not observed, are encoded in the unobserved confounder  $U$ . The causal  
189 model we selected is shown in [Figure 1](#).



*Figure 1: The graphical causal model for our research question represented as a DAG.*

190 Under this causal model, the only confounders are the vaccine strain and assay strain, and  
191 any unobserved confounders. If we assume no unmeasured confounding, then the  
192 minimal sufficient adjustment set is only the vaccine strain and assay strain. However, our  
193 goal in this project was to analyze the effect of antigenic distance as a predictor without  
194 incorporating strain-specific effects. So, we stratified our models by vaccine strain (i.e., we

195 fit all models separately for each vaccine component) and deliberately did not include a  
196 strain-specific effect.

197 To adjust for nuisance variation (potentially a source of unmeasured confounding), we  
198 included random effects to control for measurements at the same study site and on the  
199 same individual. Finally, we included specific ancestors of the outcome variable, which is  
200 not necessary to close backdoor paths and does not mitigate any biases in our estimates.  
201 However, including these variables can improve the efficiency of the estimators of interest.  
202 We included pre-vaccination titer and age specifically in our model. In our previous work,  
203 we found that sex and race/ethnicity have minimal association with the observed HAI  
204 titers, but we included them in our model for completeness. Since the majority of  
205 participants in our study were white (see the demographics table in a later section), we  
206 coded race/ethnicity as an indicator variable that was equal to 0 if the participant identified  
207 as white or Caucasian and not Hispanic or Latino, and 1 otherwise. We coded sex as 0 if  
208 the participant's sex assigned at birth was reported as male and 1 if it was reported as  
209 female. We included pre-vaccination titer in the model as-is, but since the age has a large  
210 range (from 11 to 65), we minmax scaled the age before using it in a model. Minmax scaling  
211 variables with large ranges can improve numerical stability of the model, but the model  
212 can still make predictions for any age. Similarly, we minmax scaled the birth year.

213 Finally, we note that in some models it is also possible for cross-season differences to  
214 exist when the same vaccine strain was used for multiple years in a row. I.e., we might  
215 expect post-vaccination titers to change due to repeated usage of the same vaccine.  
216 However, since some of the vaccine strains were only used for one year before being  
217 replaced, this seasonal effect is not estimable in all of our models. Therefore, we decided  
218 not to include a seasonal effect in any of the models, especially since the effect of  
219 repeated usage of the same vaccine strain was not our primary research question.

## 220 2.4 Model implementation

221 We fit two models using `brms`, a generalized additive mixed model (GAMM) and a linear  
222 mixed model (LMM). The models were identical other than the specification for the effect  
223 of antigenic distance, so we will first describe the general parts of the model. Note that in  
224 the following mathematical descriptions, we adopt bracket notation rather than subscript  
225 notation following the convention of McElreath (22) due to the large number of subscripts  
226 in our model. That is, we use the notation  $y[i]$  in place of the conventional  $y_i$ . We use  
227 subscripts to instead identify unique parameters. We also used the centered dot symbol  
228 ( $\cdot$ ) to avoid repetition when there are many valid arguments that would have the same  
229 right-hand side in a formula. For example,  $\zeta[\cdot]$  indicates that all subscripts for  $\zeta$  use the  
230 same equation.

231 We modeled our outcome (post-vaccination titer) as a Gaussian random variable, but due  
232 to the censored nature of our data we applied a censoring correction in the likelihood.  
233 Letting the outcome for a specific vaccine component be  $y$ , we assumed that

234

$$f(y[i] | \mu[i], \sigma^2) = \int_{L[i]}^{U[i]} \mathcal{N}(y[i] | \mu[i], \sigma^2) dy[i]$$

$$\sigma \sim t^+(3,0,1)$$

$$i = 1, \dots, n$$

235 where  $L[i]$  and  $U[i]$  are the lower and upper censoring bounds respectively (see the  
 236 section on censoring bounds for details),  $\mathcal{N}(\mu, \sigma^2)$  is the Gaussian (Normal) probability  
 237 density function with mean  $\mu$  and variance  $\sigma^2$ ,  $t^+(\nu, \mu, \sigma)$  is the location-scale half  
 238 Student's  $t$  distribution with degrees of freedom  $\nu$ , location parameter  $\mu$ , and scale  
 239 parameter  $\sigma$ . We chose a Student's  $t$  prior with  $\nu = 3$  degrees of freedom because the  
 240 distribution has fat tails, which allows the variance to be large if supported by the data, but  
 241 we assume *a priori* that the distribution of the variance has a finite location and scale  
 242 parameter (which is only the case when  $\nu > 2$ ). Here,  $i$  is the index for the current data  
 243 record, representing one HAI assay, and  $n$  is the total number of HAI assays (records) in the  
 244 dataset.

245 The model for the mean is shown below, including the priors for each parameter. For now,  
 246 we represent the effect of antigenic distance as a function  $g$ , which we detail with its priors  
 247 in the next section. (Note that here we refer to the  $\beta_j$  with subscripts because we treat  
 248 each as an independent parameter rather than as a vector of parameters, which would be  
 249 implied by  $\beta[j]$ . There is conceptually little difference between these notational  
 250 approaches but we feel that this notation better emphasizes the independent priors on  
 251 each  $\beta_j$ .)

252

$$\begin{aligned} \mu[i] &= \beta_0 + u[1,id[i]] + u[2,study[i]] + u[3,subtype[i]] + \\ &\quad u[4,subtype[i] \times vaccine\ strain[i]] + u[5,subtype[i] \times assay\ strain[i]] + \\ &\quad g(antigenic\ distance[i]) + \\ &\quad \beta_p(log\ pre-vaccination\ titer[i]) + \beta_a(scaled\ age[i]) + \\ &\quad \beta_y(scaled\ birth\ year[i]) + \beta_r(race/ethnicity[i]) + \beta_s(sex[i]) \\ \beta_{(.)} &\sim \mathcal{N}(0,1) \\ u[r,:] &\sim \mathcal{N}(0, \omega[r]) \quad r = 1, 2, \dots, 5 \\ \omega[r] &\sim t^+(3,0,1) \end{aligned}$$

253 The priors follow the same formulation as before, but we chose Gaussian priors for the  
 254 beta effects. Gaussian priors have flatter tails than Student's  $t$  priors, which provides a  
 255 more regularizing effect for the beta parameters – that is, we presuppose that they are  
 256 more likely to be close to zero, and our data needs to be strong enough to move the  
 257 posterior distributions away from zero before we can make any conclusions.

258 The functional form of  $g$  is the only difference between the GAMM and the LMM. In the  
 259 LMM,  $g$  takes a simple linear form:

260

$$\begin{aligned} g(antigenic\ distance[i]) &= \beta_d(antigenic\ distance[i]) \\ \beta_d &\sim \mathcal{N}(0,1) \end{aligned}$$

261 where the antigenic distance is minmax normalized for each model as described in the  
262 antigenic distance calculation section. For the GAMM, the function form of  $g$  is more  
263 complex. We modeled the antigenic distance effect using a thin-plate basis spline, which  
264 allows for the relationship to be curved in an arbitrary pattern, but constrains the fit so that  
265 rapid changes in the pattern are penalized and must be supported by data (23–27). The  
266 specific form is

267

$$g(\cdot) = \sum_{k=1}^5 \gamma[k] \cdot \phi[k](\cdot)$$
$$\gamma[k] \sim \mathcal{N}(0, \tau)$$
$$\tau \sim t^+(3, 0, 0.25)$$

268 where the  $\gamma[k]$  are coefficients which are regularized to be similar via an adaptive prior and  
269 the  $\phi[k]$  are thin-plate spline basis functions. Thin-plate splines use a low-rank  
270 approximation of the spline basis for computational efficiency, which can be tuned to  
271 balance between accuracy and efficiency. The maximal  $k$  (or size of the spline basis) we  
272 can choose is equal to the number of unique values of the predictor, so we chose  $k = 5$ ,  
273 which was estimable across all of our antigenic distance metrics. We used Student's  $t$   
274 priors for the adaptive prior on the variance of the spline coefficients so that the spline can  
275 be wiggly if supported by the data, but we chose a conservative hyperprior variance (0.25,  
276 based on a prior predictive simulation) to constrain the spline towards being flat if the  
277 signal from the data is not strong.

278 The random effects we included in the model represented sources of nuisance variation  
279 which we were interested in controlling for, but not specifically estimating. We included  
280 random effects to capture interindividual variation, variation across study sites, and direct  
281 effects of the influenza strains not explained by the antigenic distance. We included  
282 random intercepts for individuals ( $u[1, \cdot]$ ) and study sites ( $u[2, \cdot]$ ) in a typical way with  
283 regularizing priors. To control for the direct effects of influenza assay and vaccine strains,  
284 we noted that each strain was nested within a subtype, but the assay strains and vaccine  
285 strains were not themselves crossed or nested (each assay strain could appear with an  
286 arbitrary combination of different vaccine strains, although all assay strains and all  
287 vaccine strains are only ever associated with a single influenza subtype). Including the  
288 subtype effect as  $u[3, \cdot]$  and then including random effects which consider both the  
289 subtype and the vaccine strain ( $u[4, \cdot]$ ) or the assay strain ( $u[5, \cdot]$ ) allows for assay/vaccine  
290 strains within the same subtype to have a correlated effect, while assay/vaccine strains of  
291 different subtypes do not have correlated effects. Again, we assigne skeptical, regularizing  
292 priors to all of these random effects.

## 293 2.5 Censoring bounds

294 HAI titer assays, like all serial dilution assays, produce censored data values. In fact, all  
295 values produced by an HAI assay are censored. We take this censoring into account in the  
296 likelihood of our model by integrating over the censoring bounds for a given data point  $y_i$ .

297 All serial dilution assays are censored – for the case of HAI, we assume that there is some  
298 latent, true dilution  $y_i^*$  which is the minimal dilution where hemagglutination is not  
299 observed. This is likely some decimal number, and we will never observe this true value.  
300 Instead, we chose a starting dilution,  $y_{\min}$ , which is 10 in our dataset. If we observe  
301 agglutination at this starting dilution, we say the value is below the limit of detection and it  
302 is recorded as 5 in our dataset. These values are left censored. In reality, we know that the  
303 latent agglutination dilution for an assay can be any value less than 10, i.e., our censoring  
304 bounds for these assays are (0,10).

305 There is also a maximal dilution for the assay,  $y_{\max}$ , which was 20480 in our dataset. In  
306 practice, if researchers don't observe hemagglutination at any dilution, they can simply  
307 continue diluting the assay until they observe agglutination. However, a standard 96-well  
308 plate only has 12 columns, so most studies will report 20480 (the 12th serial dilution for an  
309 HAI assay starting at 10 and doubling each dilution). So these values are right censored,  
310 and the censoring bounds are  $[20480, \infty)$ . Note that the lower bound of the interval is  
311 included because the value could be exactly 20480 (though this occurs with probability  
312 zero for a continuous latent variable).

313 Finally, any other assay with a result between the limits of detection will also be interval  
314 censored, because we only observe certain dilutions. For example, if we observe inhibited  
315 hemagglutination at a dilution of 10, but agglutination occurs at a dilution of 20, we record  
316 the result as 10. However, we don't know that a dilution of 1:15 wouldn't cause inhibition,  
317 so we only know that the latent dilution is in the interval [10,20). Similarly for any value  
318  $y_{\min} < y < y_{\max}$ , the latent dilution is in the interval  $[y, 2y)$ .

319 Converting to the log scale, the censoring bounds  $L$  and  $U$  that we refer to in the previous  
320 equations are as follows (here we omit subscripts to avoid confusion with interval  
321 notation, but  $L$ ,  $U$ , and  $y$  all vary by individual while  $y_{\min}$  and  $y_{\max}$  are constant):

$$322 \quad (L, U) = \begin{cases} (-\infty, y_{\min}) & y = y_{\min} \\ [y, y + 1), & y_{\min} < y < y_{\max} \\ [y_{\max}, \infty) & y = y_{\max} \end{cases}$$

323 For our study,  $y_{\min} = \log_2(10/5) = 1$  and  $y_{\max} = \log_2(20480/5) = 12$ .

## 324 2.6 Stan implementation

325 We obtained posterior samples of the model parameters using the No U-Turn Sampler  
326 (NUTS) algorithm implemented by Stan (28,29), via the `brms` (30–32) and `cmdstanr` (33)  
327 packages for R (34). In `brms`, we specified our model formulas as

```
328 y | cens(c, y2) ~ 1 +
329   birth_year_c + age_c + sex_i + race_i +
330   log_pretiter + s(d_norm, k = 5, by = strain_type) +
331   (1 | strain_type) +
332   (1 | study) + (1 | subject_id) +
333   (1 | strain_type:vaccine_name) + (1 | strain_type:strain_name)
```

334 for the GAMMs and

```
335 y | cens(c, y2) ~ 1 +
336     birth_year_c + age_c + sex_i + race_i +
337     log_pretiter + d_norm + (1 + d_norm | strain_type) +
338     (1 | study) + (1 | subject_id) +
339     (1 | strain_type:vaccine_name) + (1 | strain_type:strain_name)
```

340 for the LMMs. We specified our prior distributions as

```
341 brms::prior(normal(0,1), class = "Intercept"),
342 brms::prior(normal(0,1), class = "b"),
343 brms::prior(student_t(3, 0, 1), class = "sd", lb = 0),
344 brms::prior(student_t(3, 0, 1), class = "sigma", lb = 0),
345 brms::prior(student_t(3, 0, 0.25), class = "sds", lb = 0)
```

346 for the GAMMS and

```
347 brms::prior(normal(0,1), class = "Intercept"),
348 brms::prior(normal(0,1), class = "b"),
349 brms::prior(student_t(3, 0, 1), class = "sd", lb = 0),
350 brms::prior(student_t(3, 0, 1), class = "sigma", lb = 0)
```

351 for the LMMs. Note that the LMMs do not have a prior for parameters of class sds, which  
352 represent the adaptive smoothing priors for spline coefficients. We sampled the models on  
353 32 chains with 200 warmup iterations and 625 post-warmup sampling iterations per chain  
354 for a total of 20,000 posterior post-warmup samples for each parameter. The effective  
355 number of samples is shown in the model diagnostics table in a later section. We also  
356 specified an adaptive delta of 0.99 and a maximum treedepth of 12, and used the  
357 recommended Stan and brms default values for all other algorithm control samples.  
358 Notably, this means each chain was initialized with a random variable – choosing smart  
359 initial values could potentially speed up the sampling, but we found that this was not  
360 sufficient to warrant further investigation for our models. The primary cause of slow  
361 sampling for our models was the large number of data points, although we also have  
362 several hierarchical parameters which can slow sampling.

## 363 2.7 Posterior marginal effects

364 To summarize our models, in Figure 2 we present a posterior contrast that we call a  
365 marginal effect, but is technically a marginal conditional effect that marginalizes some  
366 variables and is conditional on others. The posterior effect of interest is the effect of  
367 antigenic distance on post-vaccination titer, conditional on the subtype for each antigenic  
368 distance metrics (note that models were fit completely separate for each antigenic  
369 distance metric). In order to calculate this effect, we constructed counterfactual  
370 predictions to estimate for our model on an interpolated grid of antigenic distance values  
371 using the `marginalEffects` package (35).

372 The effects we present in main text Figure 2 (for each separate antigenic distance metric,  
373 which were all fit as completely separate models with the same set of effects other than  
374 the different antigenic distances) are conditional on the subtype, and represent global  
375 means of the marginal effects at the mean for the other variables included. This means  
376 that are results are conditional on the random effects in the model, but do not include  
377 random effects variances in the credible intervals for effects other than the subtype.

378 We specified interpolated counterfactual values of the normalized antigenic distance from  
379 0 to 1, spaced by 0.01, excluded the effects of random effects parameters, set categorical  
380 fixed effects values to their mode, and set continuous fixed effects (other than antigenic  
381 distance) to their mean. In this way, our marginal effects represent the expected post-  
382 vaccination titer for this population for a typical individual who is similar to those included  
383 in our study population. While using average marginal effects (AMEs) allows us to  
384 generalize our predictions to other levels of the fixed effects for a typical individual, and we  
385 could integrate out the random effects, we found that this more computationally  
386 demanding methods only served to inflate the credible interval without substantively  
387 changing the predictions we made. Since the credible interval is already quite wide, and  
388 should be interpreted conservatively as in all non-causal observational studies, we did not  
389 see the need to include nuisance variation in our predictions, since the main focus of our  
390 study was comparing the similarity of predictions across the four antigenic distance metric  
391 models, rather than specifically trying to isolate a causal effect of antigenic distance.

392 Specifically, we obtained predictions from the models on a grid defined by the following  
393 `marginaleffects` package syntax.

```
394 marginaleffects::datagrid(  
395     model = model_i,  
396     d_norm = seq(0, 1, 0.01),  
397     strain_type = c("H1N1", "H3N2", "B-Vic", "B-Yam")  
398 )
```

399 Here, `model_i` refers to each of the models that we fit.

## 400 3. Supplementary results

### 401 3.1 Annual Fluzone vaccine formulation

402 [Table 2](#) shows the strains which were included in each season's formulation of the Fluzone  
403 vaccine. We only show the formulation for the standard dose (SD) vaccine (the HD vaccine  
404 was trivalent throughout the study years we selected, while the quadrivalent formulation of  
405 the SD vaccine became available in 2015/16).

*Table 2: Strains used in the Fluzone standard dose vaccine formulation during each influenza season.*

Season	A(H1N1)	A(H3N2)	B/Victoria	B/Yamagata
2013/14	CA/09	TX/12	—	MA/12
2014/15	CA/09	TX/12	—	MA/12
2015/16	CA/09	Switz/13	Bris/08	Phu/13
2016/17	CA/09	HK/14	Bris/08	Phu/13
2017/18	MI/15	HK/14	Bris/08	Phu/13

### 406 3.2 Annual heterologous strain panel

407 The strains used in each panel are shown in [Table 3](#). A shaded cell with an X in it indicates  
 408 that the strain indicated by the current row was used as part of the HAI panel in the season  
 409 indicated by the current column.

*Table 3: Heterologous strain panel used during each influenza season.*

Subtype	Strain	2013/14	2014/15	2015/16	2016/17	2017/18
A(H1N1)	SC/18	X	X	X	X	X
	PR/34	X				
	Wei/43	X	X	X	X	X
	FM/47	X	X	X	X	X
	Den/57	X	X	X	X	X
	NJ/76	X	X	X	X	X
	USSR/77	X	X	X	X	X
	Bra/78	X			X	X
	CA/78		X	X		
	Chi/83	X	X	X	X	X
	Sing/86	X	X	X	X	X
	TX/91	X	X	X	X	X
	Bei/95	X	X	X	X	X
	NC/99	X	X	X	X	X
	SI/06	X	X	X	X	X

	Bris/07	X	X	X	X	X
	CA/09	X	X	X	X	X
	MI/15				X	X
A(H3N2)	HK/68	X	X	X	X	X
	PC/73	X	X	X	X	X
	TX/77	X	X	X	X	X
	MI/85	X	X	X	X	X
	Sich/87	X	X	X	X	X
	Shan/93	X	X	X	X	X
	Nan/95	X	X	X	X	X
	Syd/97	X	X	X	X	X
	Pan/99	X	X	X	X	X
	Fuj/02	X	X	X		
	NY/04	X	X	X	X	X
	Br/07	X				
	WI/05	X	X	X	X	X
	Uru/07		X	X	X	X
B/Pre	Per/09	X	X	X	X	X
	Vic/11	X	X	X	X	X
	TX/12	X	X	X	X	X
	Switz/13	X	X	X	X	X
	HK/14		X	X	X	X
	Sing/16					X
	Lee/40	X	X	X	X	
B/Victoria	MD/59		X	X	X	
	Sing/64		X	X	X	
B/Victoria	Vic/87				X	X

	HK/01			X	X	X
	Mal/04			X	X	X
	Vic/06			X	X	X
	Bris/08			X	X	X
	CO/17			X	X	X
B/Yamagata	Yam/88	X	X	X	X	X
	Harb/94	X	X	X	X	X
	Sich/99	X	X	X	X	X
	FL/06	X	X	X	X	X
	WI/10	X	X	X	X	X
	TX/11	X	X	X	X	X
	MA/12	X	X	X	X	X
	Phu/13	X	X	X	X	X

### 410 3.3 Strain names and abbreviations

411 Throughout the manuscript, we use abbreviated names for each strain. [Table 4](#) shows the  
 412 corresponding abbreviation for the full name of each strain.

*Table 4: Full strain names and associated abbreviations for each strain used in the study.*

Subtype	Strain name	Short name
A(H1N1)	A/H1N1/South Carolina/1/1918	SC/18
	A/H1N1/Puerto Rico/8/1934	PR/34
	A/H1N1/Weiss/1943	Wei/43
	A/H1N1/Fort Monmouth/1/1947	FM/47
	A/H1N1/Denver/1957	Den/57
	A/H1N1/New Jersey/8/1976	NJ/76
	A/H1N1/Ussr/90/1977	USSR/77
	A/H1N1/Brazil/11/1978	Bra/78
	A/H1N1/California/10/1978	CA/78
	A/H1N1/Chile/1/1983	Chi/83
	A/H1N1/Singapore/6/1986	Sing/86
	A/H1N1/Texas/36/1991	TX/91
	A/H1N1/Beijing/262/1995	Bei/95
	A/H1N1/New Caledonia/20/1999	NC/99
	A/H1N1/Solomon Islands/3/2006	SI/06
	A/H1N1/Brisbane/59/2007	Bris/07

A(H3N2)	A/H1N1/California/07/2009 A/H1N1/Michigan 45/2015 A/H3N2/Hong Kong/8/1968 A/H3N2/Port Chalmers/1/1973 A/H3N2/Texas/1/1977 A/H3N2/Mississippi/1/1985 A/H3N2/Sichuan/2/1987 A/H3N2/Shandong/9/1993 A/H3N2/Nanchang/933/1995 A/H3N2/Sydney/5/1997 A/H3N2/Panama/2007/1999 A/H3N2/Fujian/411/2002 A/H3N2/New York/55/2004 A/H3N2/Brisbane/10/2007 A/H3N2/Wisconsin/67/2005 A/H3N2/Uruguay/716/2007 A/H3N2/Perth/16/2009 A/H3N2/Victoria/361/2011 A/H3N2/Texas/50/2012 A/H3N2/Switzerland/9715293/2013 A/H3N2/Hong Kong/4801/2014 A/H3N2/Singapore/infimh-16-0019/2016	CA/09 MI/15 HK/68 PC/73 TX/77 MI/85 Sich/87 Shan/93 Nan/95 Syd/97 Pan/99 Fuj/02 NY/04 Br/07 WI/05 Uru/07 Per/09 Vic/11 TX/12 Switz/13 HK/14 Sing/16
B/Pre	B/Lee/1940 B/Maryland/1959 B/Singapore/1964	Lee/40 MD/59 Sing/64
B/Victoria	B/Victoria/02/1987 B/Hong Kong/330/2001 B/Malaysia/27127/2004 B/Victoria/326/2006 B/Brisbane/60/2008 B/Colorado/06/2017	Vic/87 HK/01 Mal/04 Vic/06 Bris/08 CO/17
B/Yamagata	B/Yamagata/16/1988 B/Harbin/7/1994 B/Sichuan/379/1999 B/Florida/4/2006 B/Wisconsin/01/2010 B/Texas/06/2011 B/Massachusetts/02/2012 B/Phuket/3073/2013	Yam/88 Harb/94 Sich/99 FL/06 WI/10 TX/11 MA/12 Phu/13

### 413 3.4 Demographic information

414 A summary of the demographic information for the individuals included in our analysis is  
 415 shown in [Table 5](#), and includes information about their reported race/ethnicity, sex  
 416 assigned at birth, age at first enrollment, and year of birth. The majority of participants  
 417 identified their race as White or Caucasian, and were assigned female at birth. All  
 418 participants from the PA and FL study sites were adults, but the UGA study site also  
 419 recruited teenagers, and all three study sites included elderly people over 65 years of age.  
 420 Most participants returned to the study site in at least one subsequent year, contributing  
 421 more than one person-year of data to the study.

**Table 5: Demographic characteristics of the study participants. Summary statistics shown are count and column percent for sex, race, and contributed person-years; and median with range for age at first enrollment, birth year, and contributed HAI assays. Demographic variables were collected by a questionnaire from participants on the date they enrolled in a study season and received a vaccine. Coding details for the demographic variables are in the Supplement.**

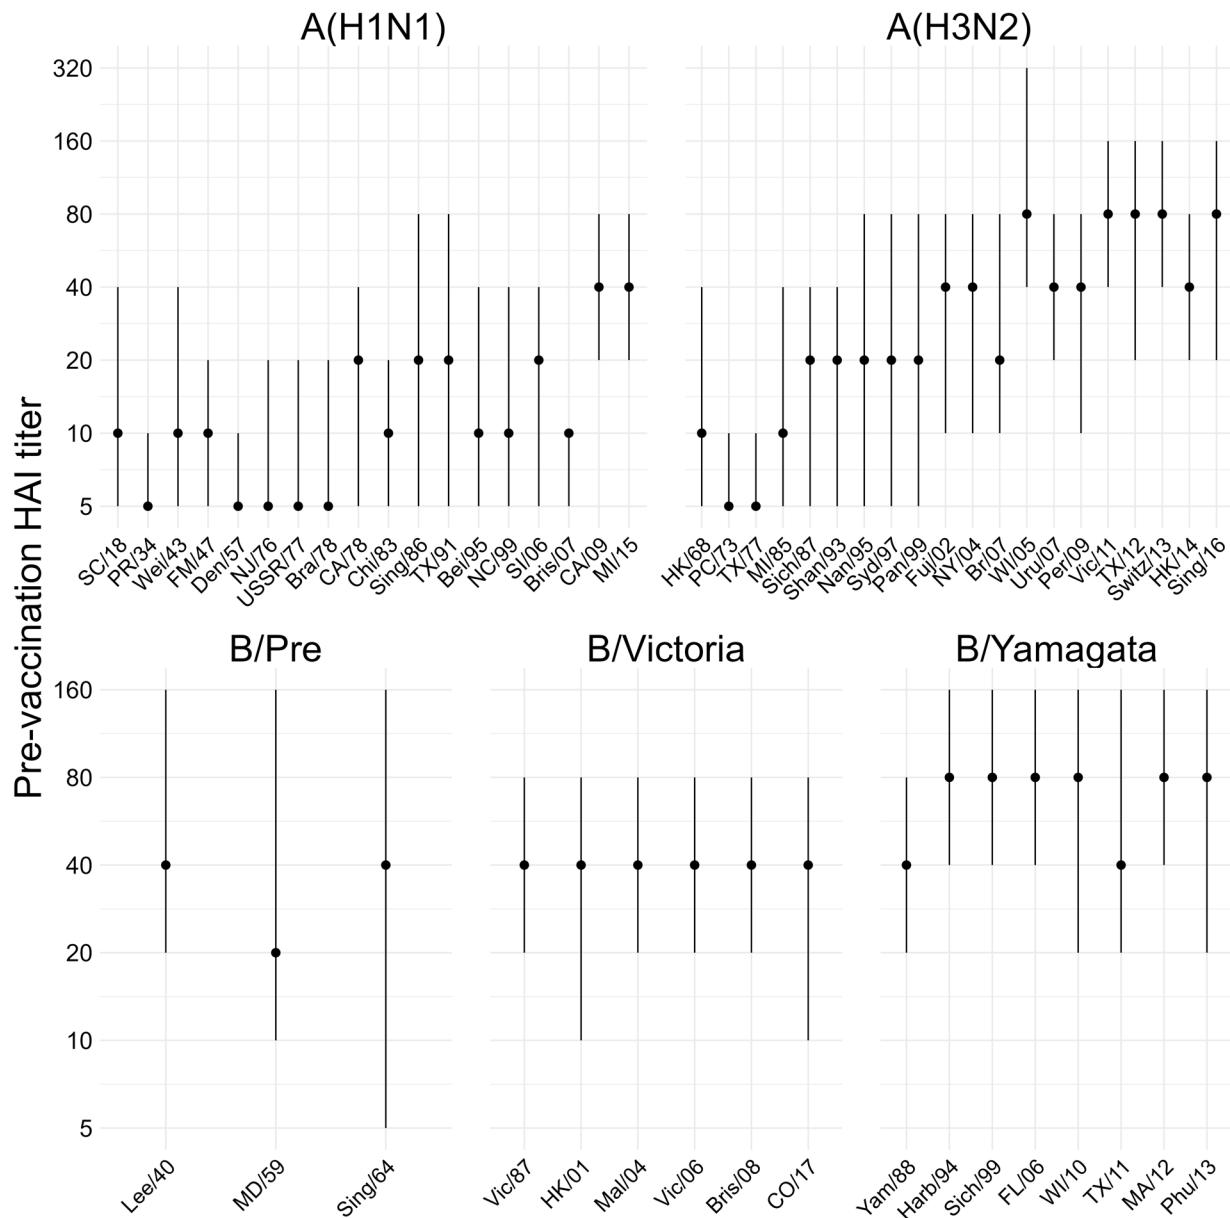
<b>Characteristic</b>	<b>FL N = 241</b>	<b>PA N = 133</b>	<b>UGA N = 303</b>	<b>Overall N = 677</b>
Sex Assigned at Birth, n (%)				
Female	184 (76)	93 (70)	168 (55)	445 (66)
Male	57 (24)	40 (30)	135 (45)	232 (34)
Race/Ethnicity, n (%)				
White	190 (79)	70 (53)	233 (77)	493 (73)
Black or African American	14 (6)	52 (39)	24 (8)	90 (13)
Other	12 (5)	8 (6)	33 (11)	53 (8)
Hispanic or Latino	24 (10)	3 (2)	13 (4)	40 (6)
Unknown	1 (0)	0 (0)	0 (0)	1 (0)
Age at First Enrollment, Median (Min - Max)	42 (20 - 80)	60 (26 - 81)	25 (12 - 83)	40 (12 - 83)
Year of Birth, Median (Min - Max)	1972 (1933 - 1996)	1954 (1932 - 1987)	1991 (1934 - 2006)	1975 (1932 - 2006)
Contributed HAI assays, Median (Min - Max)	85 (40 - 189)	94 (8 - 185)	48 (47 - 95)	52 (8 - 189)
Contributed person-years, n (%)				
1	114 (47)	44 (33)	206 (68)	364 (54)
2	52 (22)	31 (23)	97 (32)	180 (27)
3	61 (25)	32 (24)	0 (0)	93 (14)
4	14 (6)	26 (20)	0 (0)	40 (6)

422      [Figure 2](#) shows a visualization of the collected pre-vaccination titers, and [Figure 3](#) shows a  
 423      visualization of the collected post-vaccination titers, ignoring all variables except for the  
 424      assay strain.

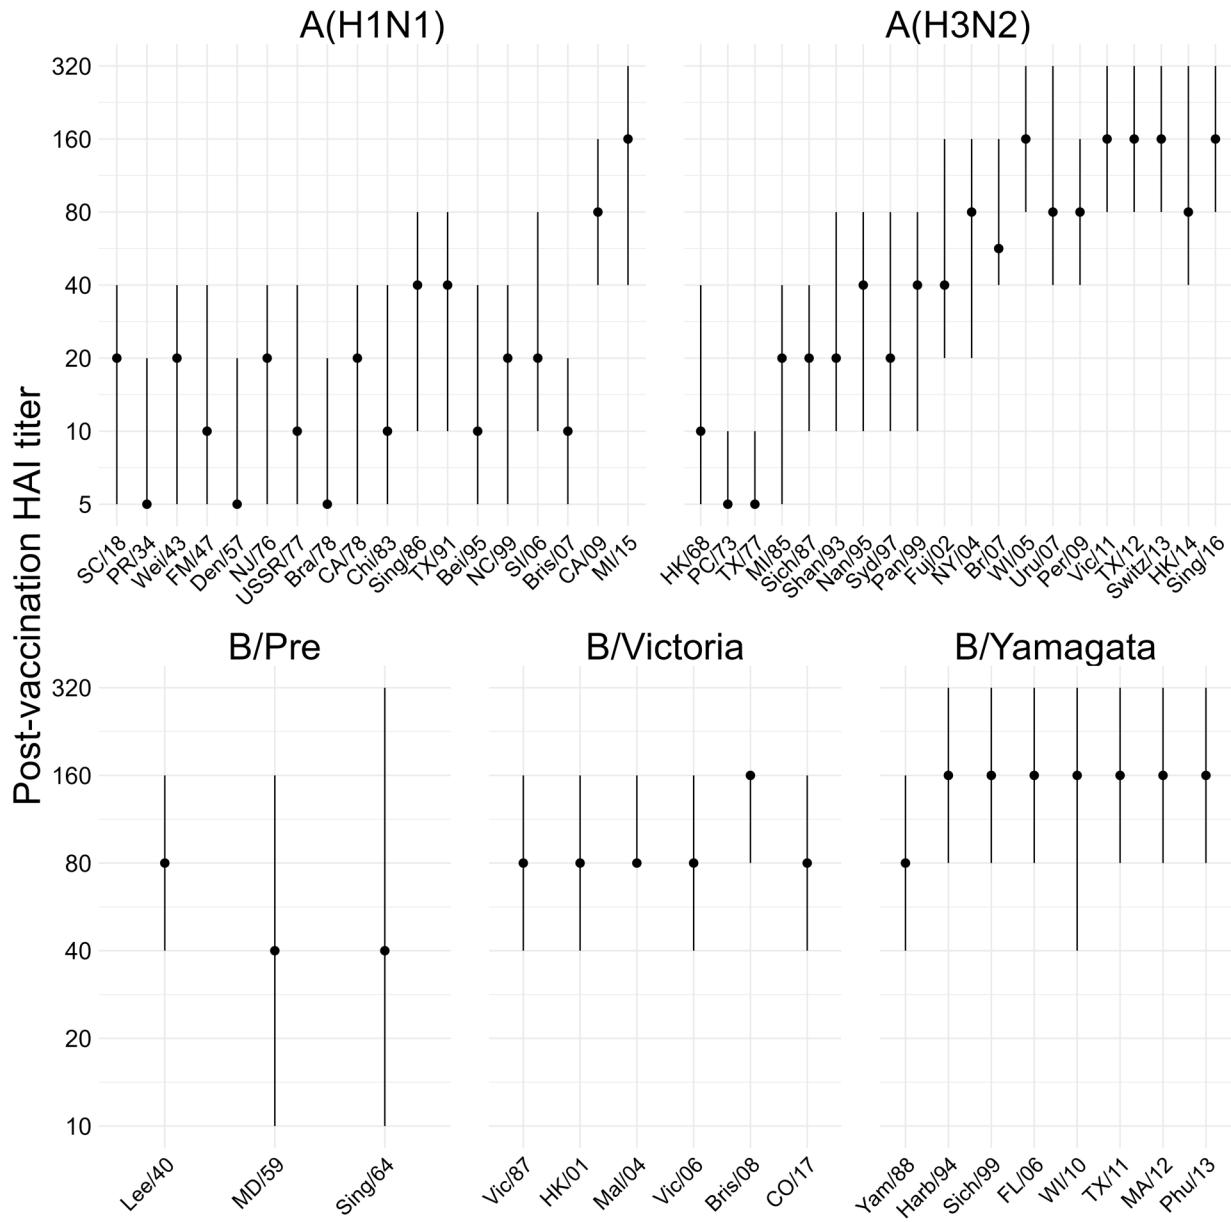
425      Qualitatively summarizing the distribution of titers to all of the assay strains from plots  
 426      alone is difficult, and the models in the main text are very helpful for understanding the  
 427      variation in post-vaccination titers. However, we can make a few observations. Most  
 428      people had some prior immunity ([Figure 2](#)) to the A(H3N2) strains which have circulated  
 429      since the 80's or 90's, with protective (40 or greater) titers to the strains from the 2000's  
 430      and onward. However, most people only had protective titers to the two most recent  
 431      A(H1N1) strains, CA/09 and MI/15 which represent the 2009 pandemic lineage. Some

432 people had immunity to older strains, but the difference was much more stark than for  
 433 A(H3N2). Many people had prior immunity to all of the B strains we examined, and the  
 434 median was 40 or greater for all of the B strains except MD/59.

435 Post-titers were, in general, higher (Figure 3). The two pandemic-like A(H1N1) strains  
 436 showed a boost on average in the population, and there was noticeable back-boosting to  
 437 some of the older A(H1N1) strains. Many of the A(H3N2) strains showed backboosting as  
 438 well, although there was not much of a response to the oldest H3N2 strains which also had  
 439 low pretiters. The median post-titers were above 40 for all of the B strains in our data, with  
 440 B/Yamagata having the highest average titers, followed by B/Victoria and then the older  
 441 (B/Pre) lineages.



*Figure 2: Pre-vaccination titers in our study to each of the assay strains. The point shows the median and the line shows the IQR.*



*Figure 3: Post-vaccination titers in our study to each of the assay strains. The point shows the median and the line shows the IQR.*

### 442 3.5 Metric agreement analysis

443 Before we built statistical models for the post-vaccination titer, we first performed a  
 444 simple unadjusted analysis of the consistency (or agreement) between the antigenic  
 445 distance measurements. As an omnibus test of agreement, we calculated the intraclass

446 correlation (ICC) across the four antigenic distance measurements, separately for each  
 447 strain type. We used a Bayesian model with a fixed effect for antigenic distance metric and  
 448 random intercepts for both assay strain and vaccine strain, and calculated the ICC as the  
 449 ratio of variance explained by the assay and vaccine strain variance components to the  
 450 total variation. The Spearman rank correlations show in the main text can be viewed as a  
 451 post-hoc analysis of the ICC which provide more information about specific comparisons.

452 Specifically, the model we fit for each subtype can be written as follows.

$$\begin{aligned}
 d[i] &\sim \mathcal{N}(\mu[i], \sigma^2) \\
 \mu[i] &= \alpha[1] \cdot I(\text{method}[i] = \text{temporal}) + \alpha[2] \cdot I(\text{method}[i] = \text{p-Epitope}) + \\
 &\quad \alpha[3] \cdot I(\text{method}[i] = \text{Grantham}) + \alpha[4] \cdot I(\text{method}[i] = \text{cartographic}) + \\
 u[1, \text{assay strain}[i]] &+ u[2, \text{vaccine strain}[i]] \\
 \alpha[k] &\sim t(3, 0.5); \quad k = 1, 2, 3, 4 \\
 u[r, \cdot] &\sim \mathcal{N}(0, \zeta[r]); \quad r = 1, 2 \\
 \zeta[r] &\sim t^+(3, 0.1) \\
 \sigma &\sim t^+(3, 0.1)
 \end{aligned}$$

453  
 454 We fit the model using Stan's NUTS sampler using 12 chains, each with 1000 warmup  
 455 iterations and 1000 post-warmup sampling iterations and an adaptive delta of 0.99. Model  
 456 diagnostics were all sufficient (data not shown, the model is easy to sample from and  
 457 samples quickly). We then calculate the ICC as

$$458 \quad \text{ICC} = \frac{\zeta[1]^2 + \zeta[2]^2}{\zeta[1]^2 + \zeta[2]^2 + \sigma^2},$$

459 over the posterior samples of all parameters. That is, the ICC represents the ratio of  
 460 variance due to strain effects only to the total variance after controlling for fixed effects. In  
 461 the psychometric literature, this is referred to as a one-way ICC for consistency – if the ICC  
 462 is close to one, it means the variance from the random effects dominates the model. We  
 463 summarized the ICC as the mean and 95% HDCl across the posterior samples.

464 As a sensitivity analysis, we considered an alternative agreement statistic based on a  
 465 different variance decomposition. We fit the same models as before, but then computed  
 466 the variance of the posterior predictions for every point in the dataset without taking the  
 467 random effects into account (the “fixed effects” predictions), i.e.

$$468 \quad \sigma_{\text{FE}}^2 = \text{Var}_{i=1}^n (\alpha[\text{method}[i]]),$$

469 where we choose the correct  $\alpha$  parameter based on the method for dataset entry  $i$  (we  
 470 omit writing all four alpha parameters and indicator functions for readability). Then, we  
 471 compute the variance of the posterior predictions for each entry in the dataset taking the  
 472 random effects and fixed effects into account:

$$473 \quad \sigma_{\text{ME}}^2 = \text{Var}_{i=1}^n (\alpha[\text{method}[i]] + u[1, \text{assay strain}[i]] + u[2, \text{vaccine strain}[i]]).$$

474 We can then compute an alternative agreement statistic as the variance ratio

475  $1 - \sigma_{FE}^2 / \sigma_{ME}^2$ ,  
 476 which will be close to one if the random effects dominate the prediction variance, or close  
 477 to zero if the random effects have only a small contribution to the prediction variance.  
 478 [Table 6](#) shows our results using this metric. All of the results indicate low agreement but  
 479 with a much higher uncertainty, and this metric is less charitable to the A(H3N2)  
 480 consistency, although we observed strong pairwise correlations between all of the  
 481 A(H3N2) metrics as shown in the main text.

*Table 6: Prediction variance ratio across all antigenic distance measurements, calculated separately for each subtype or lineage (strain type). The posterior distribution for each ratio was calculated as one minus the ratio of the prediction variance ignoring random effects to the prediction variance including random effects, estimated with a Bayesian model. Numbers shown are the mean and 95% highest density credible interval (HDCI) of the posterior distribution of variance ratios.*

Strain Type	PPD Ratio
H1N1	0.03 (-0.26, 0.32)
H3N2	0.21 (0.01, 0.39)
B-Yam	0.14 (-0.23, 0.49)
B-Vic	-0.04 (-0.75, 0.55)

### 482 [3.6 Correlation coefficients and CIs](#)

483 [Table 7](#) shows the Spearman correlation coefficients and 95% HDCIs for the correlations  
 484 shown in Figure 1 of the manuscript. The estimates and CIs are arranged by subtype in the  
 485 table in the same order in which they are shown in the plot. The coefficients may be slightly  
 486 different from the point estimates we presented in the main text due to rounding error. We  
 487 calculated the estimates and HDCIs shown here as the mean and 95% HDCI of 4000  
 488 samples from a posterior distribution created by Bayesian bootstrapping, which we  
 489 performed independently on each subtype using 4000 resamples of the observed data  
 490 points.

491 Notably, the credible intervals are quite wide for all subtypes other than A(H3N2), which  
 492 showed strong agreement in both the reliability and correlation assessments. For influenza  
 493 B subtypes, the width of the credible intervals is almost certainly due to the low number of  
 494 vaccine/assay strain pairs we observed in our dataset. For A(H1N1), we suspect that the  
 495 wide credible intervals are due to the multiple clusters in the data, which could violate the  
 496 assumptions of calculating a correlation coefficient (specifically, Spearman's rank  
 497 correlation assumes that the rank distributions are bivariate normal between the two  
 498 variables of interest, which is unlikely to hold in a variable which represents information  
 499 from multiple heterogeneous clusters). However, our point estimates which reflect low

500 overall agreement are A(H1N1) are similar to the previous results obtained by Bedford et al.  
 501 (36), and large credible intervals can indicate the lack of a strong signal in the data, so we  
 502 feel that the conclusions presented in the main text (a lack of agreement across metrics for  
 503 A(H1N1) and B subtype, and a paucity of influenza B data) are not affected by the presence  
 504 of wide credible intervals for the correlation coefficients.

*Table 7: Spearman correlation coefficients and 95% HDCl estimated by Bayesian bootstrap for each influenza subtype. Each pairwise comparison is shown only once to prevent confusion.*

Subtype		Cartographic	p-Epitope	Grantham
A(H1N1)	Temporal	0.45 ( 0.13, 0.74)	-0.10 (-0.49, 0.33)	-0.25 (-0.64, 0.17)
	Cartographi c		0.56 ( 0.29, 0.82)	0.35 (-0.01, 0.68)
	p-Epitope			0.90 ( 0.80, 0.97)
A(H3N2)	Temporal	0.93 ( 0.88, 0.97)	0.88 ( 0.80, 0.95)	0.96 ( 0.93, 0.98)
	Cartographi c		0.89 ( 0.83, 0.95)	0.93 ( 0.88, 0.97)
	p-Epitope			0.85 ( 0.72, 0.96)
B/Yamagat a	Temporal	0.66 ( 0.35, 0.89)	0.88 ( 0.76, 0.97)	0.83 ( 0.67, 0.97)
	Cartographi c		0.66 ( 0.39, 0.87)	0.67 ( 0.36, 0.91)
	p-Epitope			0.88 ( 0.73, 0.98)
B/Victoria	Temporal	0.55 ( 0.01, 0.97)	0.63 ( 0.14, 0.97)	0.71 ( 0.31, 0.99)
	Cartographi c		0.49 (-0.08, 0.95)	0.37 (-0.24, 0.88)
	p-Epitope			0.92 ( 0.73, 1.00)
Overall	Temporal	0.77 ( 0.67, 0.86)	0.63 ( 0.48, 0.75)	0.55 ( 0.39, 0.70)

---

	Cartographi c	0.82 ( 0.74, 0.89)	0.77 ( 0.68, 0.86)
	p-Epitope	0.90 ( 0.85, 0.94)	

---

### 505 3.7 Antigenic distance evenness and dispersion analysis

506 Since some of the antigenic distance metrics are more discrete than others, we calculated  
 507 the gap standard deviation as a measure of evenness of distribution across each metric.  
 508 The gap standard deviation is calculated as the standard deviation of the consecutive  
 509 differences in the sorted antigenic distance values for a given metric. That is, assume  $x$ , a  
 510 vector of measurements from  $i = 1, \dots, n$  is already sorted in increasing order so that  $x_1 \leq$   
 511  $x_2 \leq \dots \leq x_n$ . Then, the gap standard deviation is computed as

$$d_k = x_{k+1} - x_k; \quad k = 1, \dots, i-1$$

$$\bar{d} = \frac{1}{n} \sum_{k=1}^{i-1} d_i$$

$$\sigma_{\text{gap}} = \sqrt{\frac{1}{n-2} \sum_{k=1}^{i-1} (d_k - \bar{d})^2}.$$

512

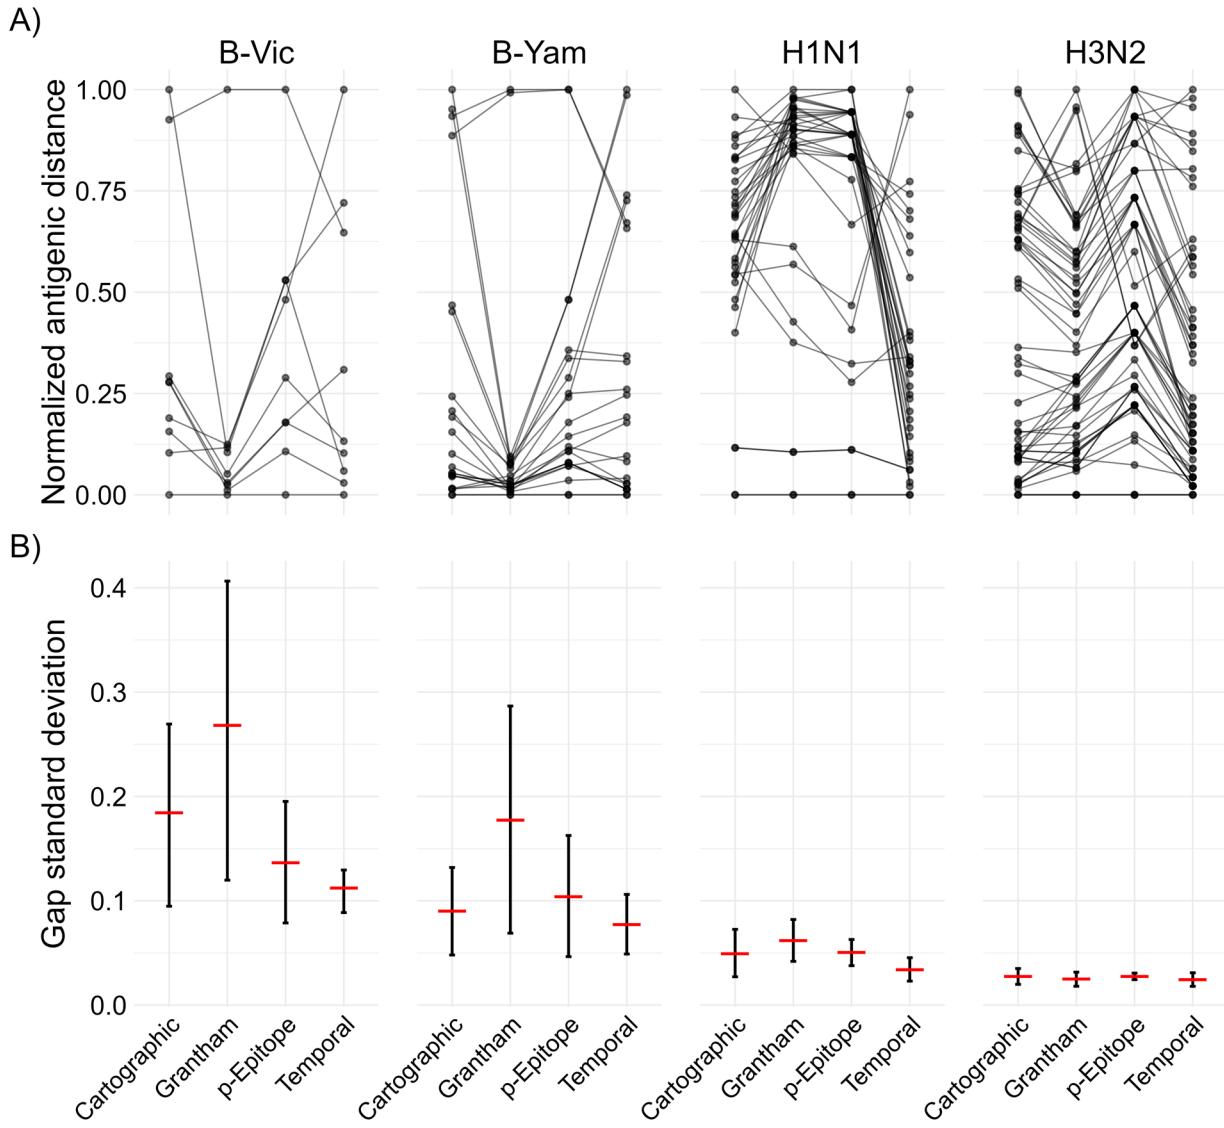
513 For a random variable with a uniform distribution,

$$514 \quad \lim_{n \rightarrow \infty} \sigma_{\text{gap}} = 0.$$

515 The different antigenic distance metrics also have different distributions in the set of  
 516 observed variables. Rather than a uniform distribution of data points across distance  
 517 space, each metric had gaps in the distribution of observed distances, which varied by  
 518 metric and subtype (Figure 4 A). The two B lineages had much larger gaps due to the  
 519 sparser historical panels. For influenza A, all metrics were more uniform for A(H3N2) than  
 520 for A(H1N1), suggesting their different evolutionary patterns across the time spanned by  
 521 the historical panel. Notably, while the temporal metric was the most uniform for all  
 522 strains (an artifact of how the historical panel was chosen), the Grantham and *p*-Epitope  
 523 metrics tend to discretize the number of potential distances and result in less uniformly  
 524 distributed values for the historical panel used in our study.

525 We quantified the uniform spread of points for each antigenic distance metric and subtype  
 526 using the gap standard deviation, where a higher gap standard deviation indicates more  
 527 irregularity in the spacing of data points. Figure 4 B shows the estimated gap standard  
 528 deviations. Both B lineages had higher gap standard deviations for all methods than either  
 529 influenza A subtype. For A(H3N2), the gap standard deviations were similar across  
 530 antigenic distance methods, and for A(H1N1) the differences were still small but larger  
 531 than A(H3N2), representing the diversity of strains in the historical panel for type A

532 influenza strains. The differences were much more noticeable for both B lineages, with  
 533 Grantham distance having notably higher gap standard deviation than the other metrics for  
 534 both influenza B lineages, indicating lower diversity in the normalized distance values.



**Figure 4: A)** Parallel coordinates plot showing how the estimated pairwise antigenic distances change for each of the antigenic distance metrics. Each line in the plot represents one vaccine strain and assay strain pair, and the connected points are the pairwise distance measured under each metric shown on the x-axis. When two lines cross, this indicates that two metrics assigned a different relative order to the pairwise combination. Note that Grantham and especially p-Epitope distances are integer-valued and concentrate measurements to specific points which potentially overlap (temporal distance is also integer valued but has enough spread to avoid a similar issue). **B)** The gap standard deviation (gap SD) for each subtype and antigenic distance metric. The posterior distribution of gap SDs was calculated using the bayesian bootstrap with

reweighting. The red horizontal bar shows the mean of the bootstrap posterior and the error bars show the 95% highest density credible interval (HDCI).

### 535 3.8 Model diagnostics

536 We examined the key model diagnostics for all of our models to ensure they converged.  
537 The main diagnostics with target criteria identified by the Stan development team (29) are:

- 538 •  $\hat{R}$ , which measures chain mixing, should be  $< 1.01$  for all parameters;  
539 • Bulk and tail ESS, measures of the number of samples drawn if all of the samples  
540 were independent, should be greater than 1000;  
541 • Number of divergent transitions should be less than 1% of samples;  
542 • Number of treedepth exceedences should be less than 1% of samples;  
543 • E-BFMI should be greater than 0.3 for all chains.

544 These diagnostics are presented in [Table 8](#).

*Table 8: Model diagnostics for the GAMMs and LMMs fit with each of the antigenic distance metrics. We show the total number of divergences out of the number of samples along with other common diagnostic criteria. For each model, we show the minimum ESS across all parameters, the minimum E-BFMI across chains, and the maximum R hat across all parameters.*

Model	Pct. Divergences	min ESS (tail)	min ESS (bulk)	min E-BFMI	max R_hat
GAMM	0.0%	5593	2574	0.48	1.02
	0.2%	6128	2831	0.51	1.02
	0.0%	2820	1483	0.45	1.03
	0.0%	4265	2103	0.52	1.02
	0.1%	3887	1914	0.52	1.02
	0.0%	3953	1870	0.48	1.02
LMM	0.0%	4293	2608	0.46	1.02
	0.0%	4325	2430	0.53	1.02
	0.3%	4286	1923	0.50	1.02
	0.0%	3967	1554	0.50	1.02
	0.0%	3903	2290	0.49	1.02
	0.0%	4953	2398	0.52	1.02

545 Most of our models had  $\hat{R}$  statistics which were 1.02 or 1.03, but these were for highly  
546 constrained parameters. Each model already takes at least 6 days to run in an HPC  
547 environment, so running the models longer is computationally infeasible due to the size of  
548 our dataset and unlikely to qualitatively change our results.

549 We also examined trace plots of the parameters to ensure there were no obvious errors  
550 (and, in general, errors in the trace plots will be noticeable in the  $\hat{R}$  statistic). We also  
551 examined the prior/posterior shrinkage and visually inspected prior/posterior plots. Since  
552 we have many models, each with thousands of parameters, we did not include the plots  
553 here. We observed good values of shrinkage (far from 1, indicating a divergence away from  
554 the prior) for most parameters, with the exception of some highly constrained parameters,  
555 typically correlations and GAMM regularizing variance parameters. Some of the random  
556 effects for individuals had poor shrinkage as well, but overall the shrinkage for random  
557 effects and for the random effects variances was far from 1. Since the GAMM was not  
558 supported by ELPD anyways, we did not investigate prior sensitivity analysis further since  
559 all of the LMM parameters had good shrinkage. Therefore, we feel safe about our choice of  
560 regularizing priors and a prior sensitivity analysis would require extensive computational  
561 time without being useful.

### 562 3.9 Prior sampling diagnostics

563 While less important for our purposes, we also sampled from the priors in order to examine  
564 the prior/posterior shrinkage and to visualize our prior predictive simulations. Such an  
565 analysis requires substantially less computational power than sampling from the posterior  
566 distribution, but we still need to ensure that we have sampled from the priors enough to get  
567 good estimates of the prior distributions of some highly constrained parameters. So, our  
568 prior sampling diagnostics are shown in [Table 9](#).

*Table 9: Model diagnostics for samples from the prior distributions for our GAMMs and LMMs. These samples are drawn only from the prior distributions and do not see the data. For each model, we show the minimum ESS across all parameters, the minimum E-BFMI across chains, and the maximum R hat across all parameters.*

Model	Pct. Divergences	min ESS (tail)	min ESS (bulk)	min E-BFMI	max R_hat
GAMM	0.0%	7920	7008	0.82	1.01
	0.0%	9218	7041	0.84	1.01
	0.0%	8789	7519	0.85	1.01
	0.0%	9100	7097	0.84	1.01
	0.0%	9013	6603	0.87	1.01
	0.0%	8805	7179	0.78	1.01
LMM	0.0%	7743	6715	0.85	1.01

---

0.0%	7412	6429	0.82	1.01
0.0%	6843	6583	0.86	1.01
0.0%	7793	6664	0.86	1.01
0.0%	7794	6586	0.81	1.01
0.0%	7618	6649	0.84	1.01

---

569 Since there are thousands of parameters per model, we do not show the shrinkage  
 570 parameters or prior distributions of all parameters here, but they are easy to produce from  
 571 the code and results we provide.

## 572 3.10 ELPD Diagnostics

573 Similar to frequentist AIC/BIC and Bayesian WAIC, ELPD relies on a computationally  
 574 efficient approximation to leave-one-out cross validation that allows estimation of a  
 575 goodness-of-fit metric without having to refit a computationally impossible number of  
 576 models. However, unlike other informatic criteria, the LOO-IC based on the leave-one-out  
 577 expected log pointwise predictive density provides diagnostics to determine if the  
 578 approximation is trustworthy (37,38). Table 10 shows the diagnostic measures for each of  
 579 our models. The maximum Pareto  $k$  diagnostic is the primary value indicating whether the  
 580 LOO-ELPD approximation is accurate – all Pareto  $k$  values (one per observation) should be  
 581 below 0.7. The  $N_{\text{eff}}$  value is the effective sample size for the approximation, and the ratio of  
 582 the effective sample size to the actual sample size should be greater than 0.5 to ensure  
 583 that the threshold of 0.7 is reliable. If the number of effective samples is greater than 2200  
 584 however, the threshold of 0.7 is useful regardless of the ratio.

*Table 10: Diagnostics for the LOO-IC ELPD approximation. Pareto  $k$  is the primary diagnostic indicating whether the approximation is trustworthy and all Pareto  $k$  values should be below 0.7. The  $N_{\text{eff}}$  is the effective sample size and  $R_{\text{eff}}$  is the ratio of the effective sample size to the true sample size – if there are too few effective samples relative to actual samples, we can get an optimistic evaluation of the approximation quality, but in general this matters less if the ESS is sufficiently high.*

---

Metric	Model	Max. Pareto k	Min. N_eff	Max. R_eff
Cartographic	GAMM	0.30	4429.2	1.00
	LMM	0.39	3522.5	1.00
p-Epitope	GAMM	0.37	4179.3	1.00
	LMM	0.37	4355.5	1.00
Grantham	GAMM	0.36	3883.5	1.00

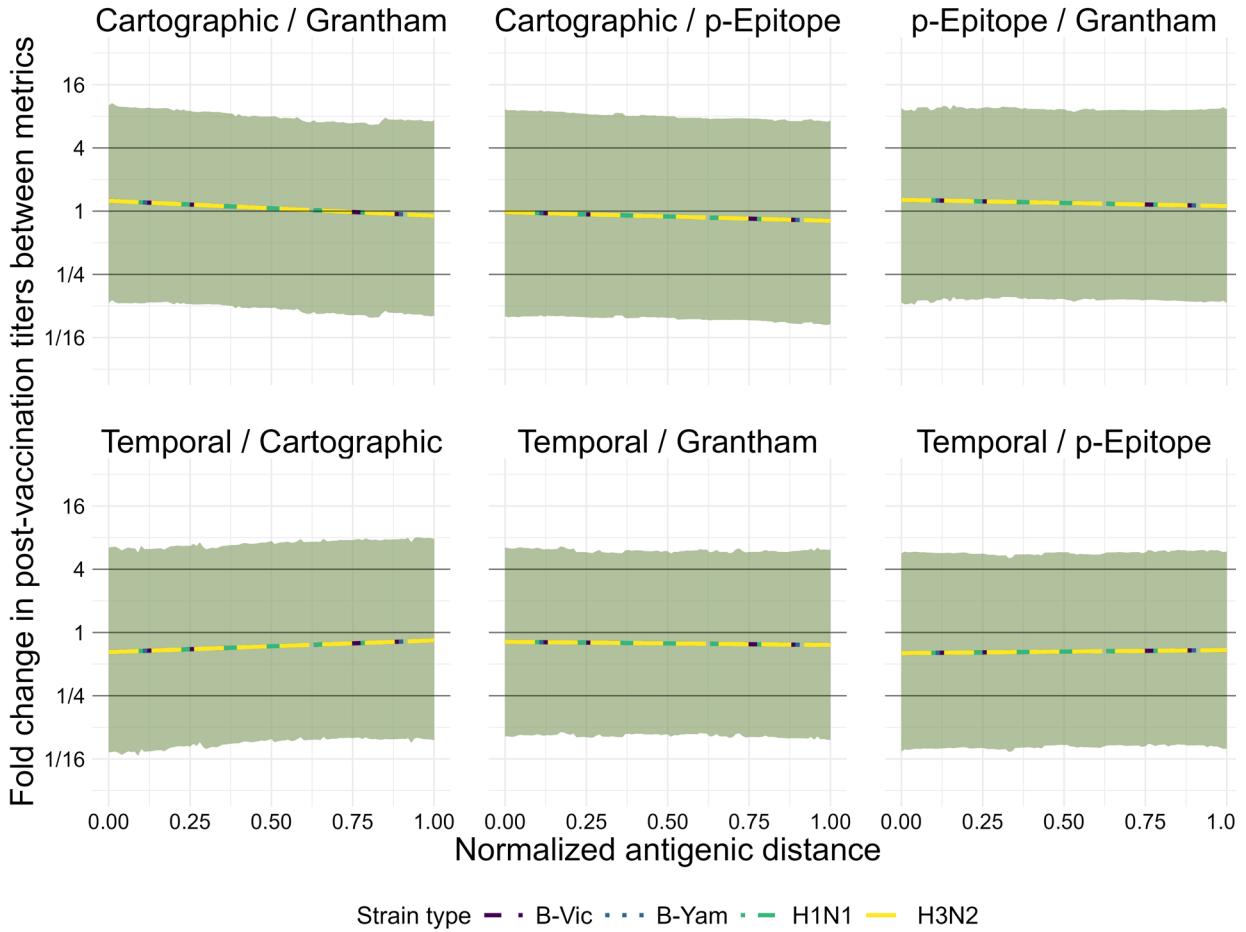
---

	LMM	0.41	3547.4	1.00
Temporal	GAMM	0.37	3788.6	1.00
	LMM	0.38	3525.1	1.00

### 585 3.11 Pointwise prediction comparisons

586 To examine the difference in predictions across each of the antigenic distance metrics, we  
 587 computed the fold change in predicted post-vaccination HAI titer conditional on  
 588 normalized antigenic distance and strain type for each unique pair of antigenic distance  
 589 metrics. We visually inspected the conditional fold changes between metrics using a limit  
 590 of agreement approach with a clinically defined threshold for whether the difference  
 591 between predictions should matter, which is commonly defined as a 4-fold change for HAI  
 592 measurements. We performed this fold change between predictions analysis for both the  
 593 GAMM and LMM with each antigenic distance metric.

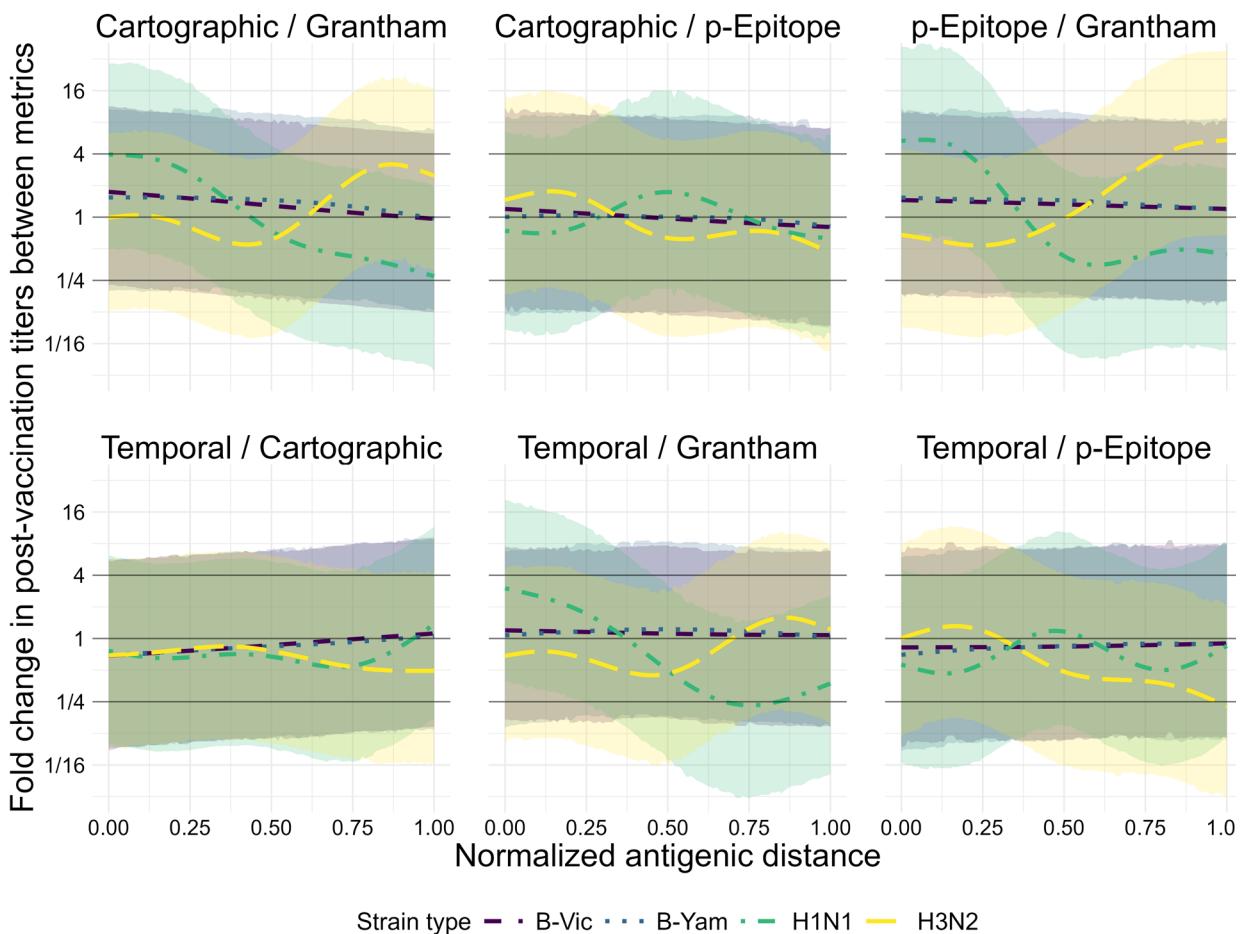
594 [Figure 5](#) shows the prediction comparisons across antigenic distance metrics for each  
 595 subtype using the LMMs. In contrast to our agreement analysis, where the A(H3N2) metrics  
 596 showed the strongest agreement across metrics (and the highest pairwise correlations),  
 597 A(H3N2) was the only subtype with noticeable trends in the contrasts between metrics. In  
 598 particular, all of the comparisons with *p*-Epitope for A(H3N2) had a noticeable trend – even  
 599 though the mean fold change in predictions always stayed within the measurement error  
 600 boundaries we set *a priori*, sometimes the credible interval did not fully cover the  
 601 measurement error boundaries and there was a noticeable slope. These trends suggested  
 602 that *p*-Epitope measurements underestimated the expected change in post-vaccination  
 603 titer compared to Grantham and cartographic distance, while *p*-Epitope overestimated the  
 604 difference compared to temporal methods. These results suggest that perhaps  
 605 biochemical features like glycosylation sites or changes to the virus outside of the  
 606 immunodominant epitope region are important, because these features are detected by  
 607 cartographic and Grantham distance, but not by *p*-Epitope distance.



*Figure 5: Pairwise comparisons of predictions (from the LMMs) between each unique set of two metrics. The y-axis shows the fold change in predictive titers between metrics, and the two metrics being compared in each subplot are shown as the subplot labels. Each line represents the predictions for the first metric in the pair at a given antigenic distance value divided by the predictions for the second metric in the pair. Color and linetype correspond to different strain types. The solid black lines on the plot are reference lines at a value of 1 for no effect, and at 4 and 1/4, effect values which would represent a clinically notable deviation in HAI predictions beyond what is expected from measurement error. Lines represent the mean of the posterior distribution of the contrast and the colored ribbons represent the 95% highest density credible interval (HDCI) for each strain type in each subplot.*

608 **Figure 6** shows the prediction comparisons across antigenic distance metrics for each  
 609 subtype using the GAMMs. Even though the GAMM was not supported by our ELPD  
 610 analysis, we used the GAMM for analyzing pairwise differences in predictions in case the  
 611 nonlinear signal was biologically important with a weak signal. Unlike our simple  
 612 correlation analysis, this analysis examines the predicted protection for an average  
 613 individual exposed to an antigenically distant strain after vaccination, rather than only

614 taking antigenic distance into account. We saw that the fold change in predicted HAI titers  
 615 was almost always less than four for every pairwise comparison between two metrics. A  
 616 four-fold change in HAI titer is considered a clinically relevant difference between two  
 617 measurements, so in almost every case we saw that changing the antigenic distance  
 618 metric would not lead to a clinically relevant difference in predicted post-vaccination HAI  
 619 titer. The primary exception was strain type A(H1N1), which exceed 40 at a few antigenic  
 620 distance values for some of the pairwise comparisons (around a normalized antigenic  
 621 distance of 0.25 for the cartographic/Grantham and Cartographic/p-Epitope comparisons,  
 622 and around a normalized antigenic distance of 0.75 for the Grantham/temporal distance  
 623 comparisons). Due to the large standard errors and the number of comparisons we make,  
 624 we are comfortable attributing these fluctuations to measurement error, although the large  
 625 variability across antigenic clusters for A(H1N1) strains (pdm-like vs. non-pdm-like) could  
 626 contribute as well.



*Figure 6: Pairwise comparisons of predictions (from the GAMMs) between each unique set of two metrics. The y-axis shows the fold change in predictive titers between metrics, and the two metrics being compared in each subplot are shown as the subplot labels. Each line represents the predictions for the first metric in the pair at a given antigenic distance value divided by the predictions for the second metric in the pair. Color and*

*linetype correspond to different strain types. The solid black lines on the plot are reference lines at a value of 1 for no effect, and at 4 and 1/4, effect values which would represent a clinically notable deviation in HAI predictions beyond what is expected from measurement error. Lines represent the mean of the posterior distribution of the contrast and the colored ribbons represent the 95% highest density credible interval (HDCI) for each strain type in each subplot.*

627 However, the differences in comparisons for A(H3N2) was not completely trivial either.  
628 [Figure 6](#) shows that for A(H3N2), the temporal distance overwhelming underestimates the  
629 fold change in predictions for the largest antigenic distances compared to both Grantham  
630 and  $p$ -Epitope measurements, with some interesting trends in the comparisons between  
631 cartographic distance as well. These results support our conclusion that further research  
632 into which of these metrics actually captures useful and interesting features is warranted,  
633 because it is difficult to tell whether we are capturing noise from our study or actual  
634 patterns that suggest different metrics are identifying different relevant characteristics of  
635 the viruses.

636 In both models, nearly all contrast predictions fall within the clinically irrelevant reference  
637 bounds, although the credible intervals for all predictions are wide because our bayesian  
638 models fairly account for many sources of uncertainty in the data. However, our results for  
639 the GAMM model suggest some interesting exceptions for the A(H1N1) strains that are  
640 likely related to the pandemic-like and non-pandemic-like cluster differences. Our results  
641 for the GAMM and LMM model for A(H3N2) seem to suggest that perhaps different metrics  
642 are picking up different relevant features, as we noted in the main text discussion.

643 3.12 Vaccine-specific Predictions

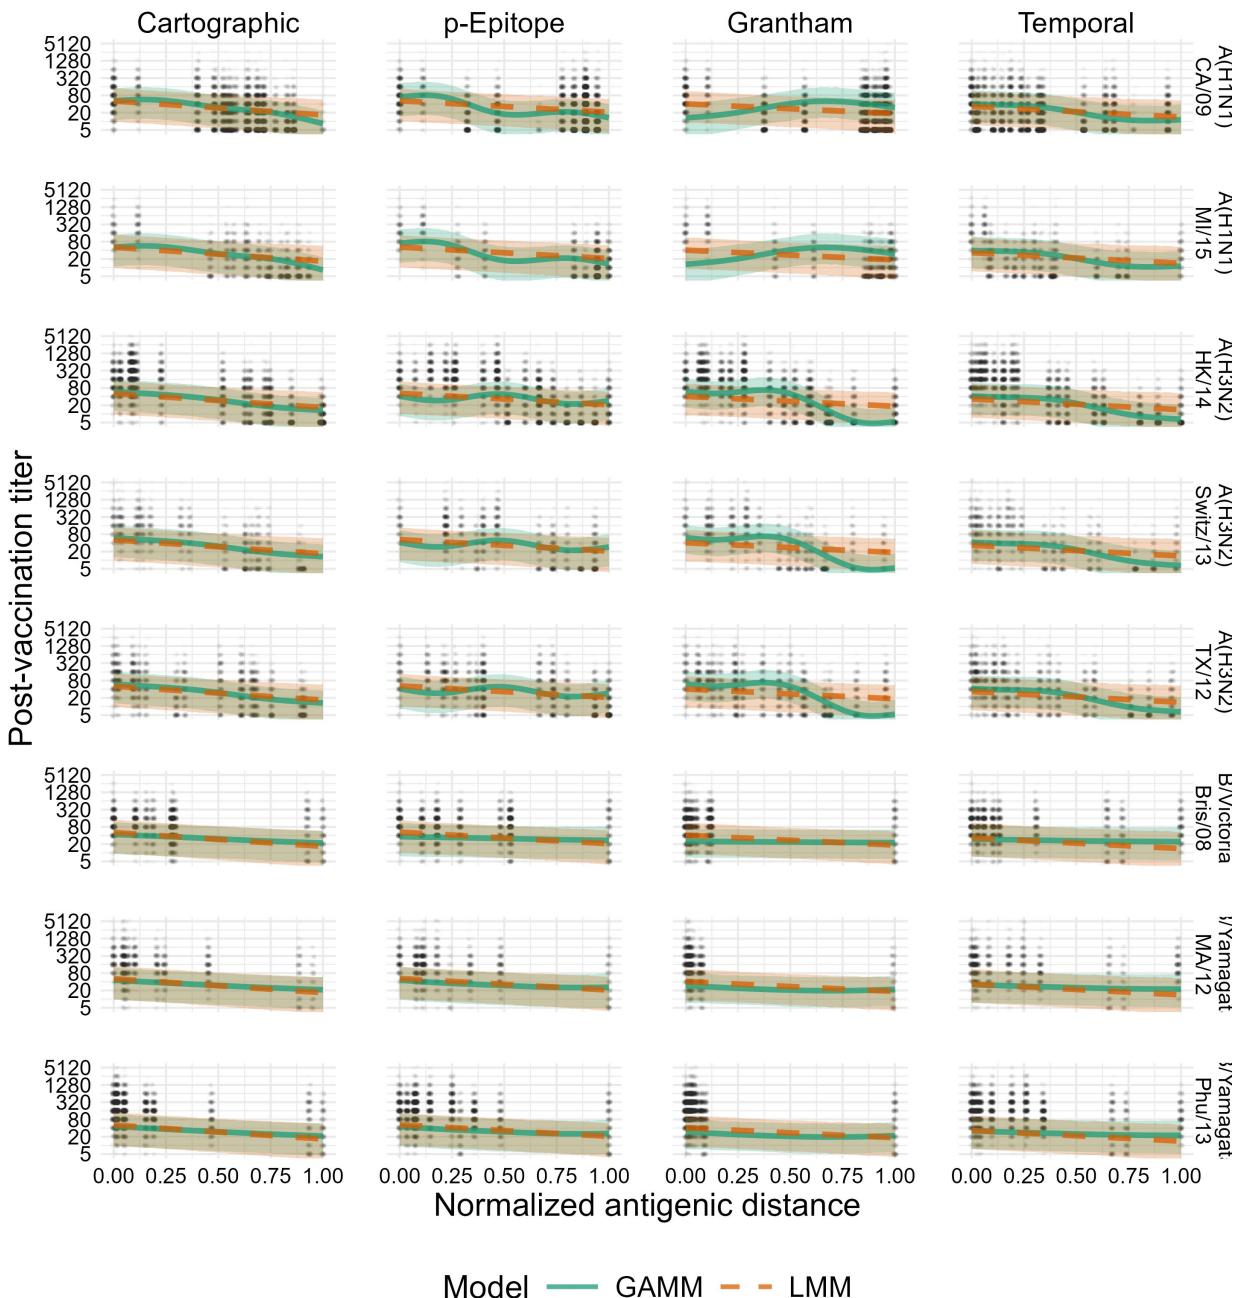


Figure 7: Model predictions for both the GAMM and LMM, conditional on the vaccine strain rather than only on the subtype (shown in the main text). Solid green lines and green ribbons show the mean and 95% highest density continuous interval (HDCI) for GAMM predictions. Dashed orange lines and orange ribbons show the mean and 95% HDCI for LMM predictions. Circular points show the data values. Each subplot shows the model predictions for a particular subtype (changes by row) and distance metric

(changes by column). Outcomes shown on the plot are predicted post-vaccination titers for an average individual to an average strain.

644 In order to analyze the differences between vaccine strains, we also examined the results  
645 conditional on the specific strain used in a vaccine (for a given subtype). Within a given  
646 subtype, there were no striking results across the different vaccine components. In the  
647 main text, we show that direct causal effects of vaccine and assay strain contribute very  
648 little to the variance in the outcomes after controlling for antigenic distance.

### 649 3.13 Alternative distance metrics

650 Because we chose to include the dominant *p*-Epitope distance and Grantham's distance  
651 in our final manuscript, we also analyzed other sequence-based and biochemical  
652 distances to determine if our arbitrary choice was misleading and we should consider  
653 further antigenic distance measures. So, we compared the dominant *p*-Epitope distance  
654 (2,39) with the *p*-all-Epitope distance (39,40) and the Hamming distance (41). Furthermore,  
655 we compared Grantham's distance (4) with the Hamming distance and with the FLU  
656 substitution model, an evolutionary amino acid substitution matrix model derived  
657 specifically for influenza sequences (42) (there are other indices like Grantham's index, but  
658 we felt that the comparison to a model specifically for influenza amino acid substitutions  
659 was sufficient).

660 When we compared the FLU substitution model to Grantham's distance and the simple  
661 Hamming distance, we found that all three metrics were highly correlated for all subtypes  
662 except A(H1N1), with relatively small credible intervals from Bayesian bootstrapping  
663 ([Table 11](#)). For A(H1N1), Grantham and Hamming distances were also highly correlated,  
664 but the correlations between Hamming and FLU substitution and Grantham and FLU  
665 substitution distances were moderate at best, with credible intervals that covered quite  
666 low values. The original study which developed the FLU substitution matrix used a mix of  
667 influenza virus sequences across multiple proteins and types/subtypes (42), so it is  
668 unclear why the difference would be so stark for A(H1N1). Regardless, because the  
669 difference was only noticeable for A(H1N1) we decided to use the Grantham distance in  
670 our main analysis. Despite high similarity to the Hamming distance across all subtypes,  
671 Grantham distance contains more information by design and better antigenic coverage of  
672 Influenza B strains in a future study might reveal further differences between Grantham  
673 distance and Hamming distance.

*Table 11: Pairwise Spearman rank correlations between antigenic distance values using the Grantham, FLU Substitution, and Hamming distance metrics. We calculated correlations between two distances using the normalized distance values between every vaccine/assay strain pair for the given subtype. Numbers shown are the mean and 95% highest density continuous interval (HDCI) calculated by Bayesian bootstrapping.*

Subtype		Grantham	FLU Substitution
A(H1N1)	Hamming	0.91 ( 0.82, 0.98)	0.62 ( 0.35, 0.85)
	Grantham		0.41 ( 0.05, 0.75)
A(H3N2)	Hamming	0.99 ( 0.98, 1.00)	0.97 ( 0.94, 0.99)
	Grantham		0.96 ( 0.93, 0.98)
B/Yamagata	Hamming	0.98 ( 0.93, 1.00)	0.94 ( 0.86, 1.00)
	Grantham		0.91 ( 0.80, 0.98)
B/Victoria	Hamming	0.99 ( 0.97, 1.00)	0.97 ( 0.89, 1.00)
	Grantham		0.97 ( 0.88, 1.00)
Overall	Hamming	0.99 ( 0.97, 1.00)	0.94 ( 0.91, 0.96)
	Grantham		0.91 ( 0.88, 0.95)

674 We also examined the pairwise Spearman correlations between the Hamming distance,  
 675 the (dominant) *p*-Epitope method which we present in the main analysis, and the *p*-all-  
 676 Epitope distance, which is calculated by averaging the Hamming distance across all 5 of  
 677 the immunodominant HA epitope sites. Again, the correlations were overall high [Table 12](#),  
 678 with A(H1N1) displaying a notably lower correlation across differences. These  
 679 supplementary results suggest that different antigenic distance metrics may have the  
 680 strongest effect on understanding the immune response to A(H1N1), probably in  
 681 accounting for notable differences across clusters. The multiple clusters in A(H1N1)  
 682 antigens are the primary differentiating factor from the ladder-like continuously  
 683 evolutionary pattern in A(H3N2) and might explain the differences, although we lack the  
 684 ability to analyze this further.

*Table 12: Pairwise Spearman rank correlations between antigenic distance values using the Grantham, FLU Substitution, and Hamming distance metrics. We calculated correlations between two distances using the normalized distance values between every vaccine/assay strain pair for the given subtype. Numbers shown are the mean and 95% highest density continuous interval (HDCI) calculated by Bayesian bootstrapping.*

Subtype		<i>p</i> -Epitope	<i>p</i> -All-Epitope
A(H1N1)	Hamming	0.93 ( 0.87, 0.98)	0.89 ( 0.79, 0.98)
	<i>p</i> -Epitope		0.89 ( 0.79, 0.97)
A(H3N2)	Hamming	0.85 ( 0.71, 0.97)	0.92 ( 0.83, 0.99)

	p-Epitope		0.97 ( 0.94, 0.99)
B/Yamagata	Hamming	0.92 ( 0.82, 0.99)	0.96 ( 0.90, 0.99)
	p-Epitope		0.97 ( 0.91, 1.00)
B/Victoria	Hamming	0.94 ( 0.80, 1.00)	0.99 ( 0.97, 1.00)
	p-Epitope		0.95 ( 0.84, 1.00)
Overall	Hamming	0.91 ( 0.86, 0.95)	0.95 ( 0.92, 0.98)
	p-Epitope		0.97 ( 0.95, 0.98)

685 Due to the relative consistency across these other antigenic distance metrics, we did not  
 686 fit further models to other antigenic distance metrics. The models require a great deal of  
 687 computational time and power, and since we found overall good agreement between these  
 688 additional metrics (and a large amount of disagreement within A(H1N1), as we found for  
 689 the metrics in the main analysis), we felt that this did not justify a further investigation.  
 690 However, a future study with an expansive panel of serological data to A(H1N1) and  
 691 A(H3N2) strains to further explore why A(H1N1) metrics have lower agreement would be  
 692 useful for further understanding these results.

## 693 4. References

- 694 1. Landau WM. The targets r package: A dynamic make-like function-oriented pipeline  
 695 toolkit for reproducibility and high-performance computing. *Journal of Open Source  
 696 Software* [electronic article]. 2021;6(57):2959. (<https://doi.org/10.21105/joss.02959>)
- 697 2. Gupta V, Earl DJ, Deem MW. **Quantifying influenza vaccine efficacy and antigenic  
 698 distance**. *Vaccine*. 2006;24(18):3881–3888.
- 699 3. Pan Y, Deem MW. **Prediction of influenza B vaccine effectiveness from sequence  
 700 data**. *Vaccine*. 2016;34(38):4610–4617.
- 701 4. Grantham R. **Amino Acid Difference Formula to Help Explain Protein Evolution**.  
 702 *Science*. 1974;185(4154):862–864.
- 703 5. Wilks S. Racmacs: Antigenic cartography macros. 2024.
- 704 6. Arhami O, Rohani P. **Topolow: A mapping algorithm for antigenic cross-reactivity  
 705 and binding affinity assays**. *Bioinformatics*. 2025;btaf372.
- 706 7. National Library of Medicine (US), National Center for Biotechnology Information.  
 707 GenBank. 1982;
- 708 8. Clark K, Karsch-Mizrachi I, Lipman DJ, et al. **GenBank**. *Nucleic Acids Research*.  
 709 2016;44(D1):D67–D72.

- 710 9. The UniProt Consortium. [UniProt: The Universal Protein Knowledgebase in 2025](#).  
711 *Nucleic Acids Research*. 2025;53(D1):D609–D617.
- 712 10. Shu Y, McCauley J. [GISAID: Global initiative on sharing all influenza data – from](#)  
713 [vision to reality](#). *Eurosurveillance*. 2017;22(13):30494.
- 714 11. Elbe S, Buckland-Merrett G. [Data, disease and diplomacy: GISAID's innovative](#)  
715 [contribution to global health](#). *Global Challenges*. 2017;1(1):33–46.
- 716 12. Reid AH, Fanning TG, Hultin JV, et al. [Origin and evolution of the 1918 "Spanish"](#)  
717 [influenza virus hemagglutinin gene](#). *Proceedings of the National Academy of Sciences of*  
718 *the United States of America*. 1999;96(4):1651–1656.
- 719 13. Smeenk CA, Brown EG. [The influenza virus variant A/FM/1/47-MA possesses single](#)  
720 [amino acid replacements in the hemagglutinin, controlling virulence, and in the matrix](#)  
721 [protein, controlling virulence as well as growth](#). *Journal of Virology*. 1994;68(1):530–534.
- 722 14. Lee M-S, Yang C-F. [Cross-reactive H1N1 antibody responses to a live attenuated](#)  
723 [influenza vaccine in children: Implication for selection of vaccine strains](#). *The Journal of*  
724 *Infectious Diseases*. 2003;188(9):1362–1366.
- 725 15. Barman S, Krylov PS, Fabrizio TP, et al. [Pathogenicity and transmissibility of North](#)  
726 [American triple reassortant swine influenza A viruses in ferrets](#). *PLoS pathogens*.  
727 2012;8(7):e1002791.
- 728 16. Garten RJ, Davis CT, Russell CA, et al. [Antigenic and genetic characteristics of](#)  
729 [swine-origin 2009 A\(H1N1\) influenza viruses circulating in humans](#). *Science (New York,*  
730 *N.Y.)*. 2009;325(5937):197–201.
- 731 17. Bauer K, Schrader C, Suess J, et al. [Neuraminidase inhibitor susceptibility of](#)  
732 [porcine H3N2 influenza A viruses isolated in Germany between 1982 and 1999](#). *Antiviral*  
733 *Research*. 2007;75(3):219–226.
- 734 18. Xu X, Kilbourne ED, Hall HE, et al. [Nonimmunoselected Intrastrain Genetic](#)  
735 [Variation Detected In Pairs Of Highyielding Influenza A \(H3N2\) Vaccine And Parental](#)  
736 [Viruses](#). *Journal of Infectious Diseases*. 1994;170(6):1432–1434.
- 737 19. Mohr PG, Deng Y-M, McKimm-Breschkin JL. [The neuraminidases of MDCK grown](#)  
738 [human influenza A\(H3N2\) viruses isolated since 1994 can demonstrate receptor binding](#).  
739 *Virology Journal*. 2015;12:67.
- 740 20. Krystal M, Young JF, Palese P, et al. [Sequential mutations in hemagglutinins of](#)  
741 [influenza B virus isolates: Definition of antigenic domains](#). *Proceedings of the National*  
742 *Academy of Sciences of the United States of America*. 1983;80(14):4527–4531.
- 743 21. Yamashita M, Krystal M, Fitch WM, et al. [Influenza B virus evolution: Co-circulating](#)  
744 [lineages and comparison of evolutionary pattern with those of influenza A and C viruses](#).  
745 *Virology*. 1988;163(1):112–122.

- 746 22. McElreath R. *Statistical Rethinking: A Bayesian Course with Examples in R and Stan*.  
747 2nd ed. Boca Raton: Taylor and Francis, CRC Press; 2020.
- 748 23. Wood SN. *Generalized Additive Models: An Introduction with R, Second Edition*. 2nd  
749 ed. New York: Chapman and Hall/CRC; 2017.
- 750 24. Wood SN. Stable and efficient multiple smoothing parameter estimation for  
751 generalized additive models. *Journal of the American Statistical Association*.  
752 2004;99(467):673–686.
- 753 25. Wood SN. Fast stable restricted maximum likelihood and marginal likelihood  
754 estimation of semiparametric generalized linear models. *Journal of the Royal Statistical  
755 Society (B)*. 2011;73(1):3–36.
- 756 26. Wood SN, N., Pya, et al. Smoothing parameter and model selection for general  
757 smooth models (with discussion). *Journal of the American Statistical Association*.  
758 2016;111:1548–1575.
- 759 27. Wood SN. *Generalized additive models: An introduction with R*. 2nd ed. Chapman  
760 and Hall/CRC; 2017.
- 761 28. Carpenter B, Gelman A, Hoffman MD, et al. *Stan: A Probabilistic Programming  
762 Language*. *Journal of Statistical Software*. 2017;76(1):1–32.
- 763 29. Stan Development Team. *Stan Modeling Language Users Guide and Reference  
764 Manual*. 2024;
- 765 30. Bürkner P-C. *brms: An R package for Bayesian multilevel models using Stan*. *Journal  
766 of Statistical Software*. 2017;80(1):1–28.
- 767 31. Bürkner P-C. *Advanced Bayesian multilevel modeling with the R package brms*. *The  
768 R Journal*. 2018;10(1):395–411.
- 769 32. Bürkner P-C. *Bayesian item response modeling in R with brms and Stan*. *Journal of  
770 Statistical Software*. 2021;100(5):1–54.
- 771 33. Gabry J, Češnovar R, Johnson A, et al. Cmdstanr: R interface to 'CmdStan'.  
772 2024.(<https://github.com/stan-dev/cmdstanr>)
- 773 34. R Core Team. *R: A language and environment for statistical computing*. Vienna,  
774 Austria: R Foundation for Statistical Computing; 2024.(<https://www.R-project.org/>)
- 775 35. Arel-Bundock V, Greifer N, Heiss A. *How to interpret statistical models using  
776 marginaleffects for R and Python*. *Journal of Statistical Software*. 2024;111(9):1–32.
- 777 36. Bedford T, Suchard MA, Lemey P, et al. *Integrating influenza antigenic dynamics  
778 with molecular evolution*. *eLife*. 2014;3:e01914.

- 779 37. Vehtari A, Gelman A, Gabry J. Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. *Statistics and Computing*. 2017;27(5):1413–1432.
- 780
- 781 38. Vehtari A, Simpson D, Gelman A, et al. Pareto Smoothed Importance Sampling. 2024;(<https://arxiv.org/abs/1507.02646>). (Accessed April 24, 2025)
- 782
- 783 39. Pan K, Subieta KC, Deem MW. A novel sequence-based antigenic distance measure for H1N1, with application to vaccine effectiveness and the selection of vaccine strains. *Protein Engineering Design and Selection*. 2011;24(3):291–299.
- 784
- 785
- 786 40. Anderson CS, McCall PR, Stern HA, et al. Antigenic cartography of H1N1 influenza viruses using sequence-based antigenic distance calculation. *BMC bioinformatics*. 2018;19(1):51.
- 787
- 788
- 789 41. Hamming RW. Error detecting and error correcting codes. *The Bell System Technical Journal*. 1950;29(2):147–160.
- 790
- 791 42. Dang CC, Le QS, Gascuel O, et al. FLU, an amino acid substitution model for influenza proteins. *BMC Evolutionary Biology*. 2010;10(1):99.
- 792